



OPEN
DAYLIGHT
S U M M I T





How to Design OpenDaylight Applications for Best Scale and Performance

Jan Medved, Distinguished Engineer, Cisco

Tony Tkacik, Architect, Cisco

#ODSummit

Agenda

- Architectural Overview & Background
- General Recommendations
- MD-SAL:
 - The Data Store
 - Messaging: RPCs and Notifications
- Coming in Beryllium:
 - MD-SAL APIs
 - Clustering
- Q&A



#ODSummit

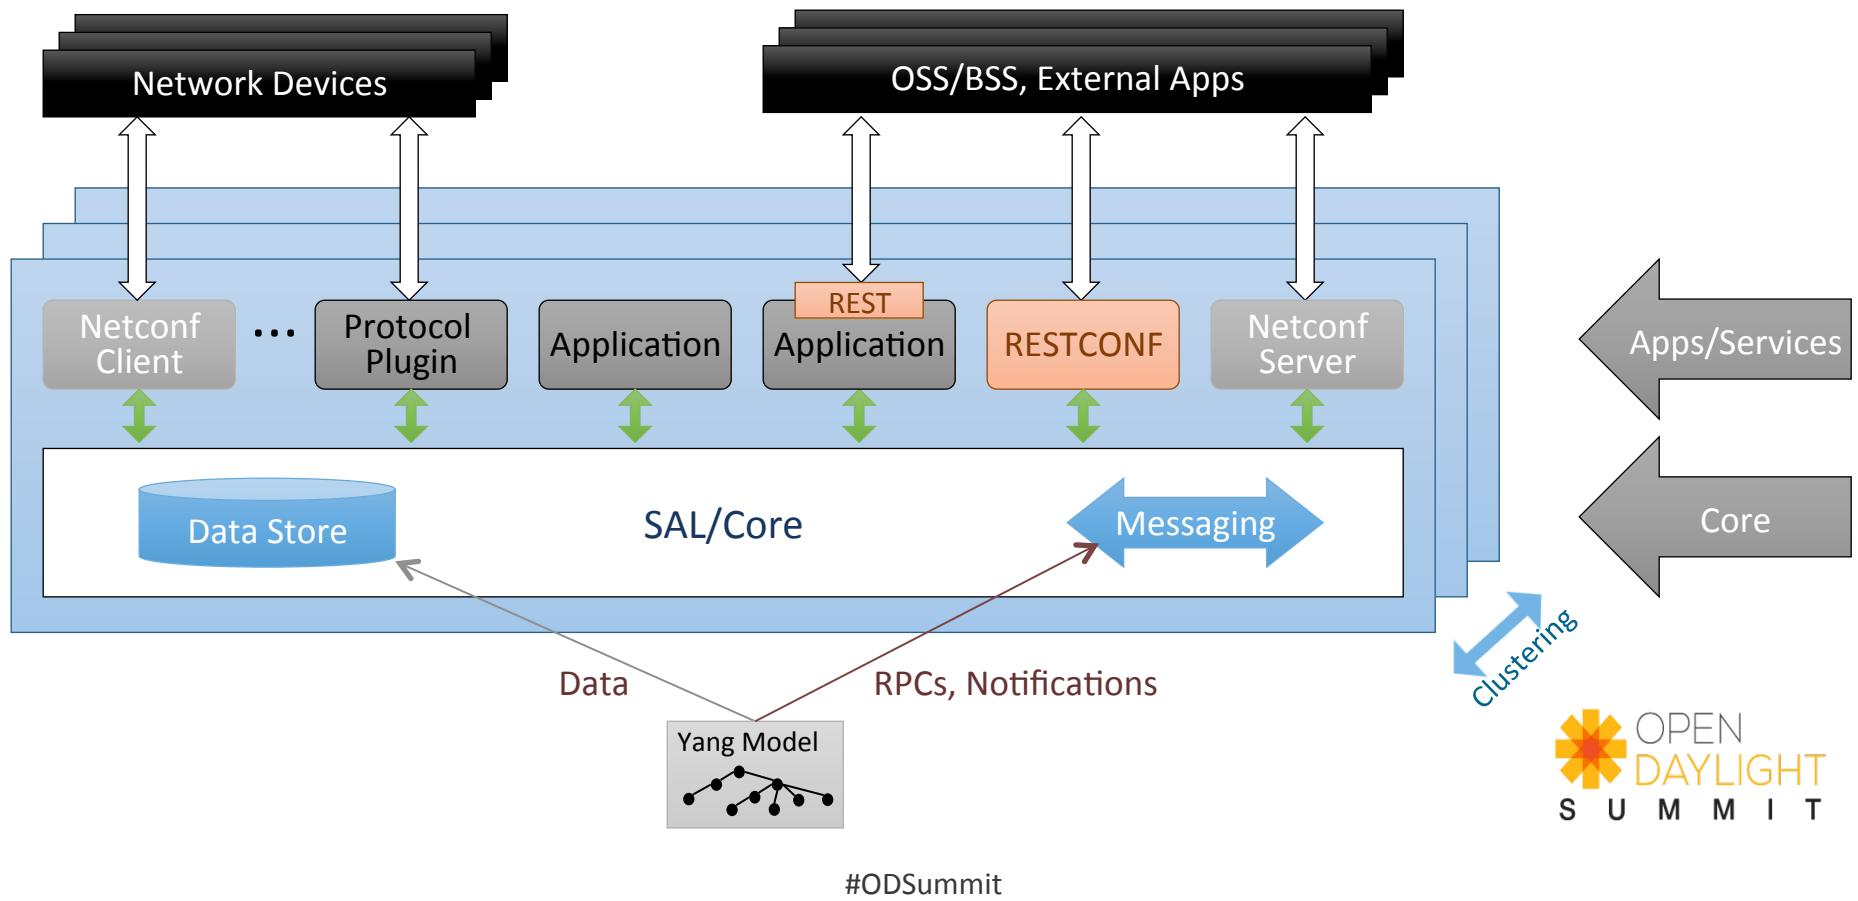


Architectural Overview & Background

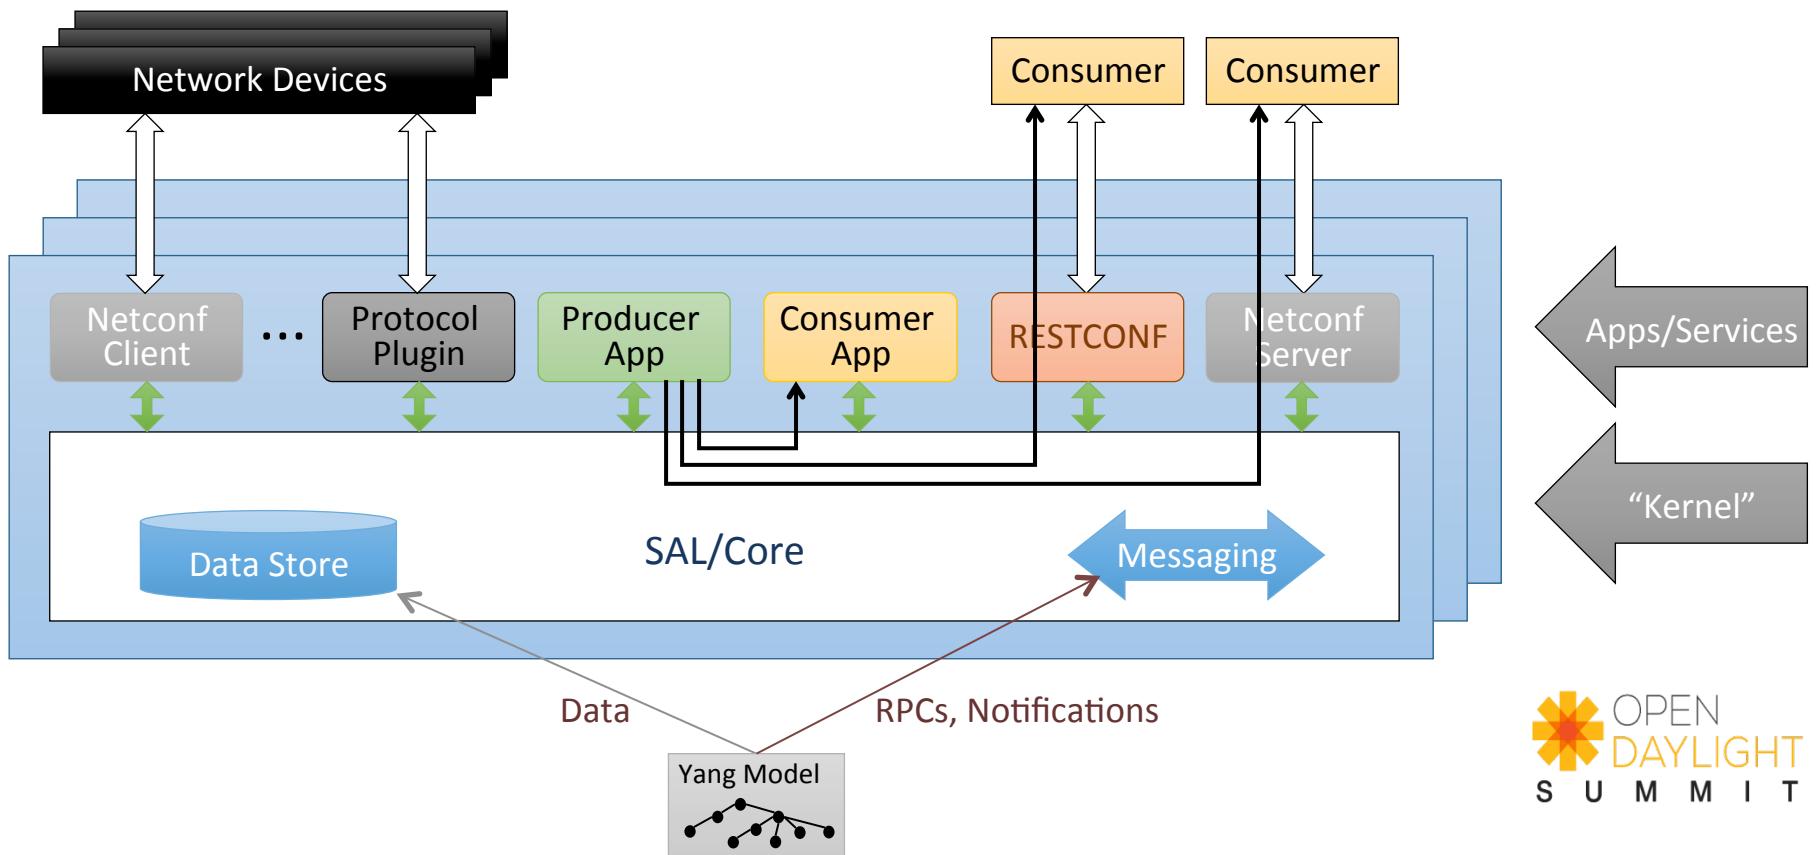
#ODSummit



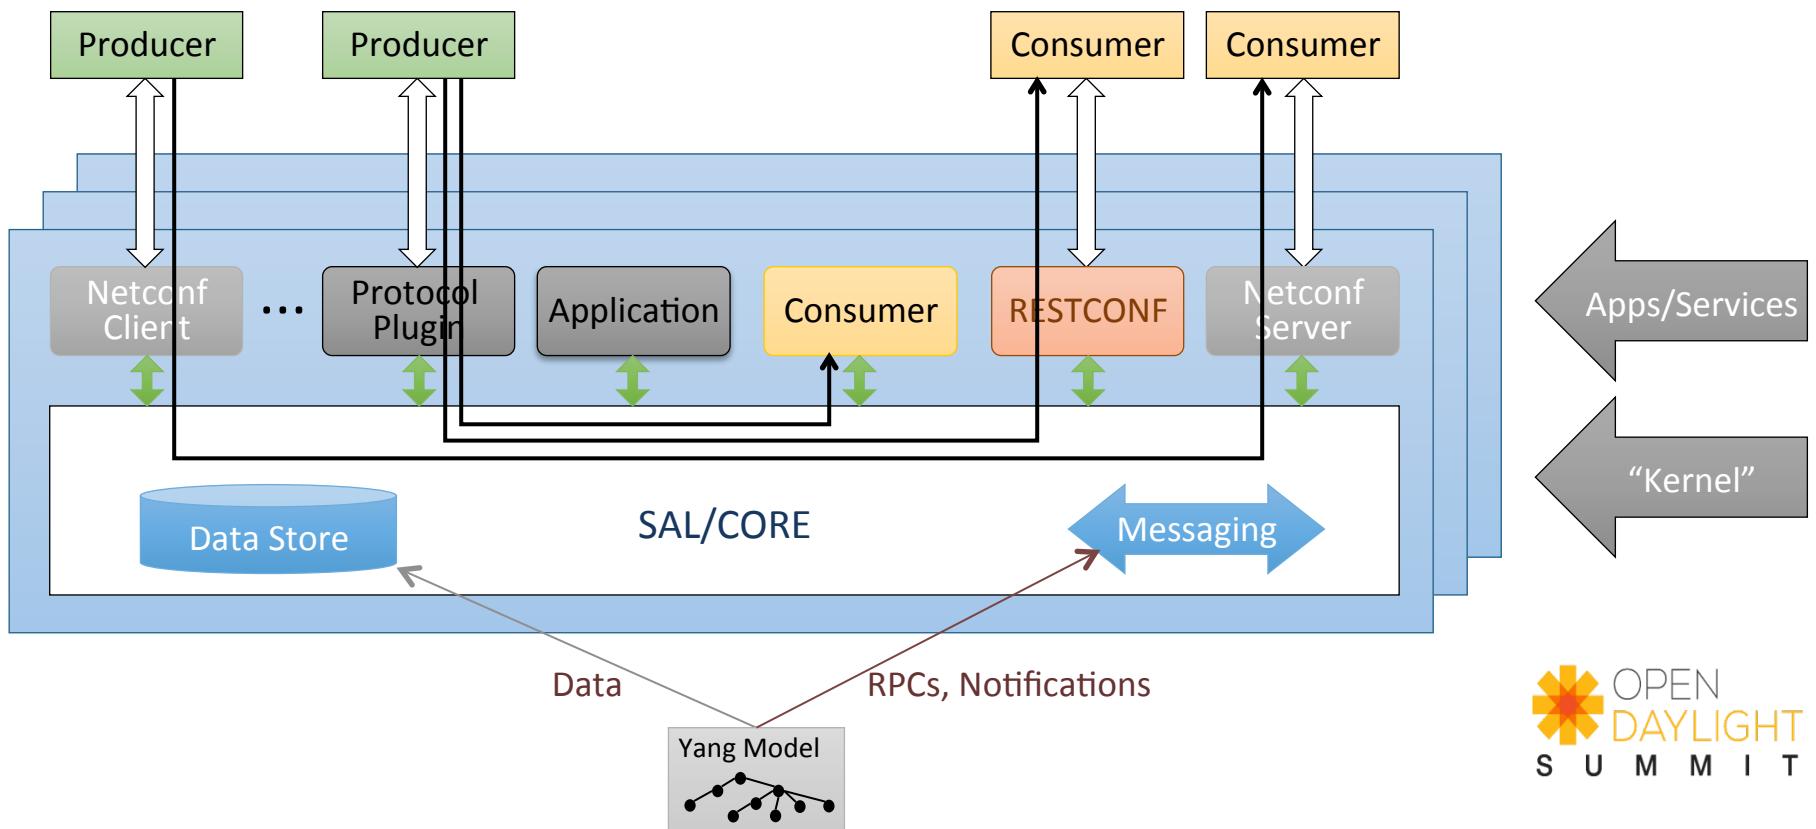
ODL Software Architecture



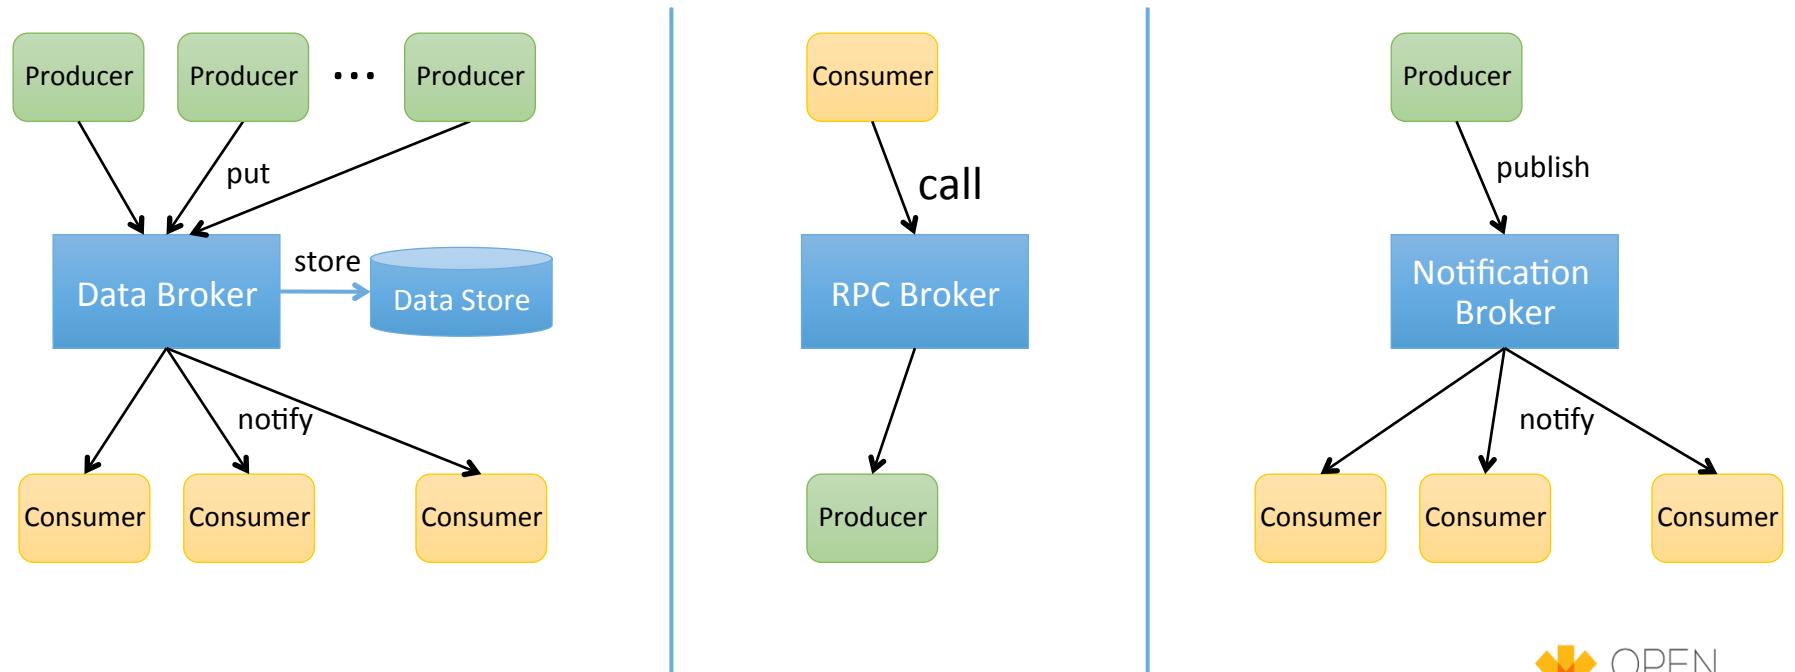
MD-SAL: Connecting Producers & Consumers



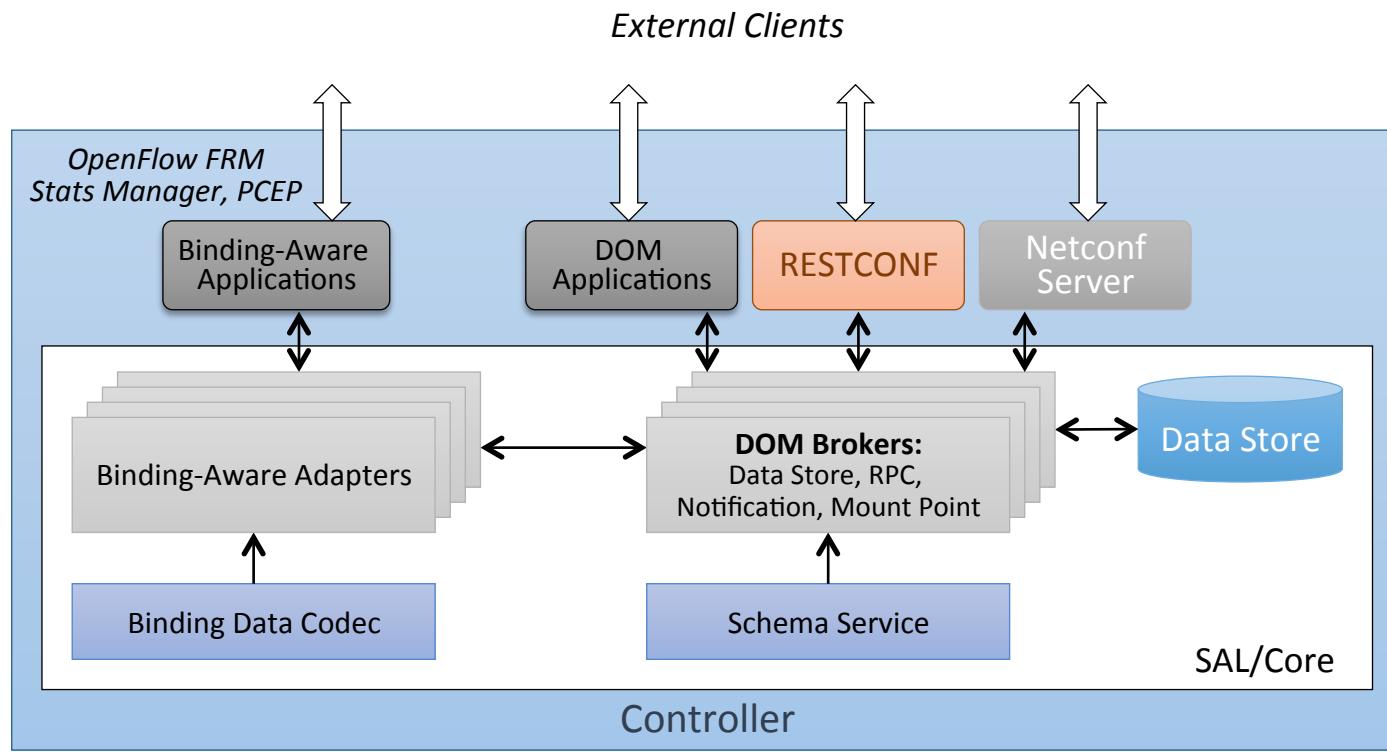
MD-SAL: Connecting Producers & Consumers



MD-SAL: The Brokers



MD-SAL: The Internals



#ODSummit



General Recommendations

#ODSummit



Using Java Futures correctly

- Avoid blocking on Futures unless necessary
- Register a callback invoked once the operation is completed:
 - Use `Futures.addCallback()`
- Assume that code returning Java Futures runs in a different thread and will take some time to complete



#ODSummit

Other Recommendations

- Programming:
 - Logging:
 - Incorrect use of logging may decrease performance:
 - Use `LOG.isDebugEnabled()` if log statement computation is CPU-heavy
 - Logging best practices:
https://wiki.opendaylight.org/view/BestPractices/Logging_Best_Practices
 - Avoid unnecessary synchronization
- Environment tuning:
 - Garbage collection
 - TCP stack configuration



#ODSummit



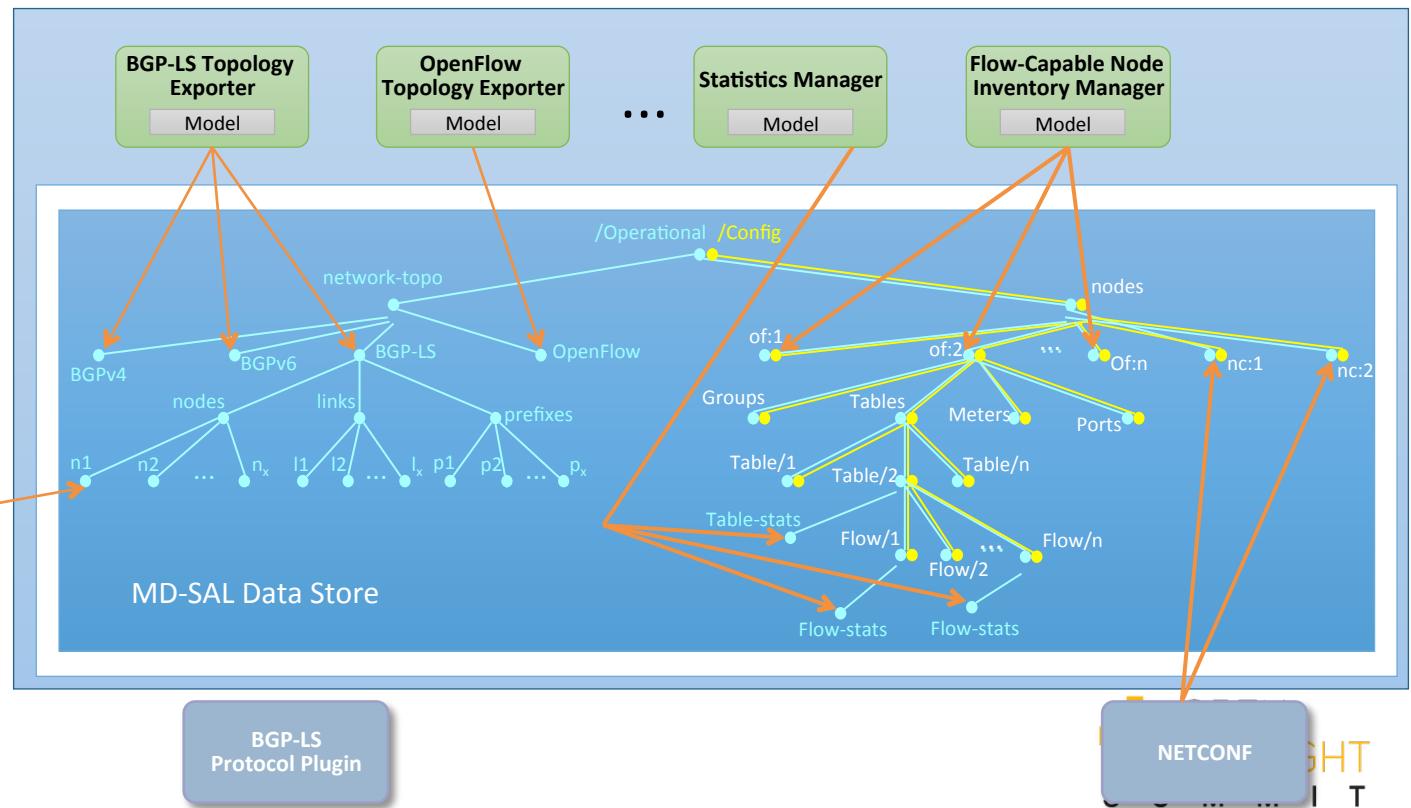
MD-SAL: The Data Store

#ODSummit



Data Store Key Concepts

- Yang data is a tree
- Two Logical Data Stores:
 - Config
 - Operational
- Unified View
- InstanceIdentifier:
 - Pointer to a node



#ODSummit

Data Store Recommendations

- Design proper structure of your data (depth vs width)
- Use **DataTreeChangeListener**
 - instead of repetitive READs
 - instead DataChangeListener
- Use **WriteOnlyTransaction** over **ReadWriteTransaction**
- Use **TransactionChain** if your transactions depends on each other
- Where possible, use **PUT** over **MERGE**.
- Use „**pingpong-broker**“ for best write batching performance with **WriteOnlyTransactions**
- **Single writer** for any given subtree



#ODSummit

Data Store Performance Considerations

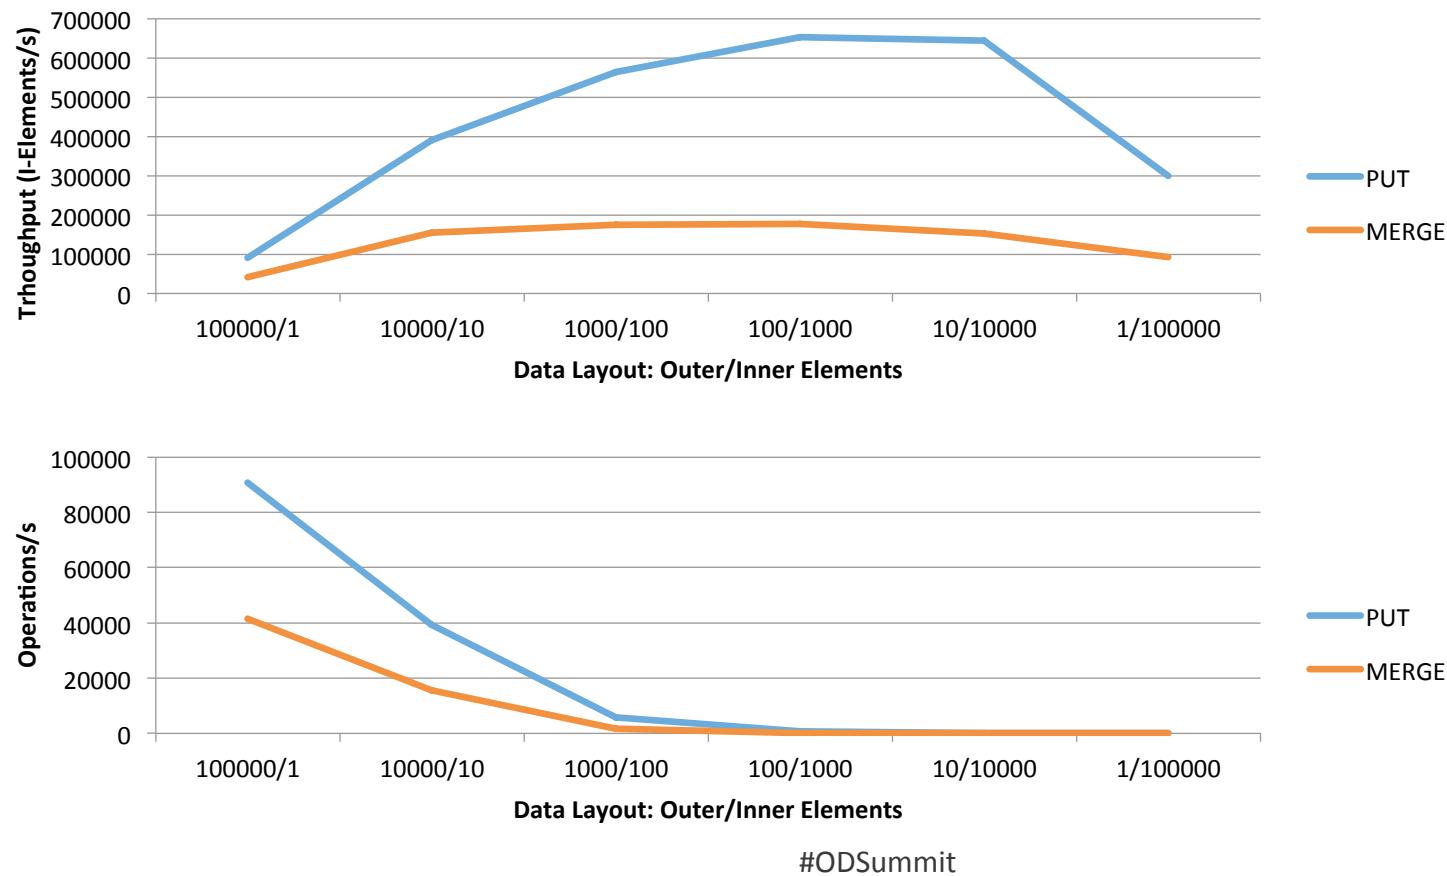
- Structure of data: breadth vs. depth
- Merge vs. Put vs. Delete
- Persistence: On/Off
- Data Broker API: Binding-Aware vs. DOM
- Simple Transaction vs. Transaction Chain
- Number of operations per submit
- Test Environment:
 - 2014 Mac:
 - 16 Gig 1600 MHz DDR3 RAM / 2.2 GHz Intel Core i7 / OSX 10.10.3 / SSD
 - Java 8 / 4G heap / G1GC
 - No clustering, single node operation



#ODSummit

Data Store Benchmark: Data Layout

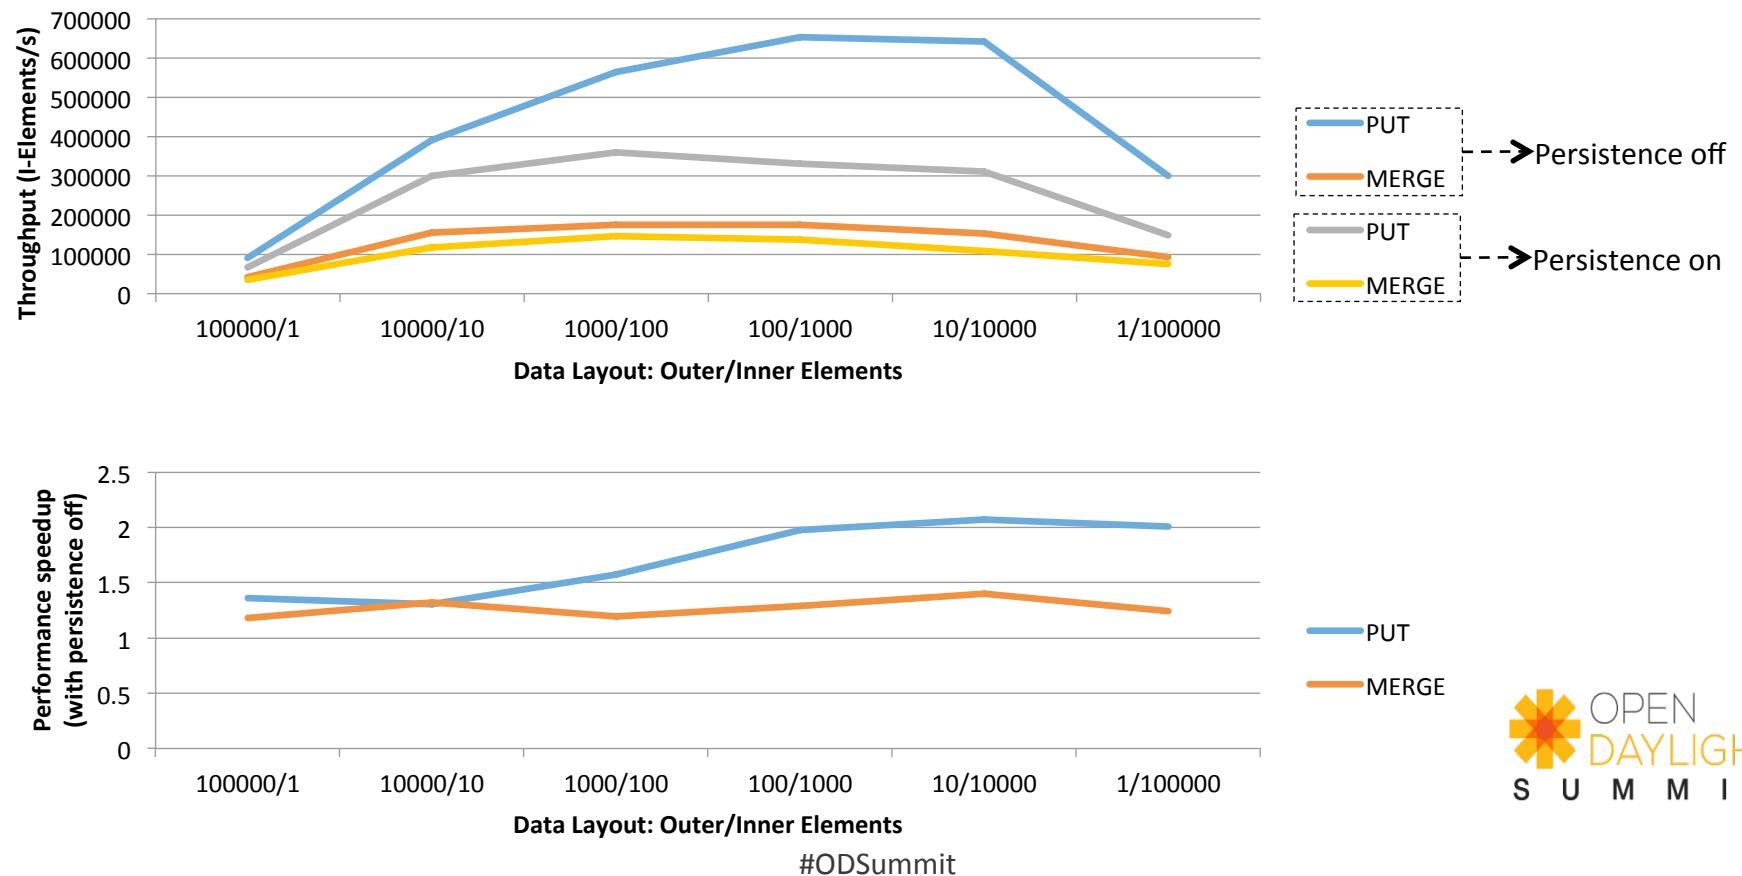
Depth vs. breadth: Data structure impact (Persistence off, transaction chain, pingpong buffer)



#ODSummit

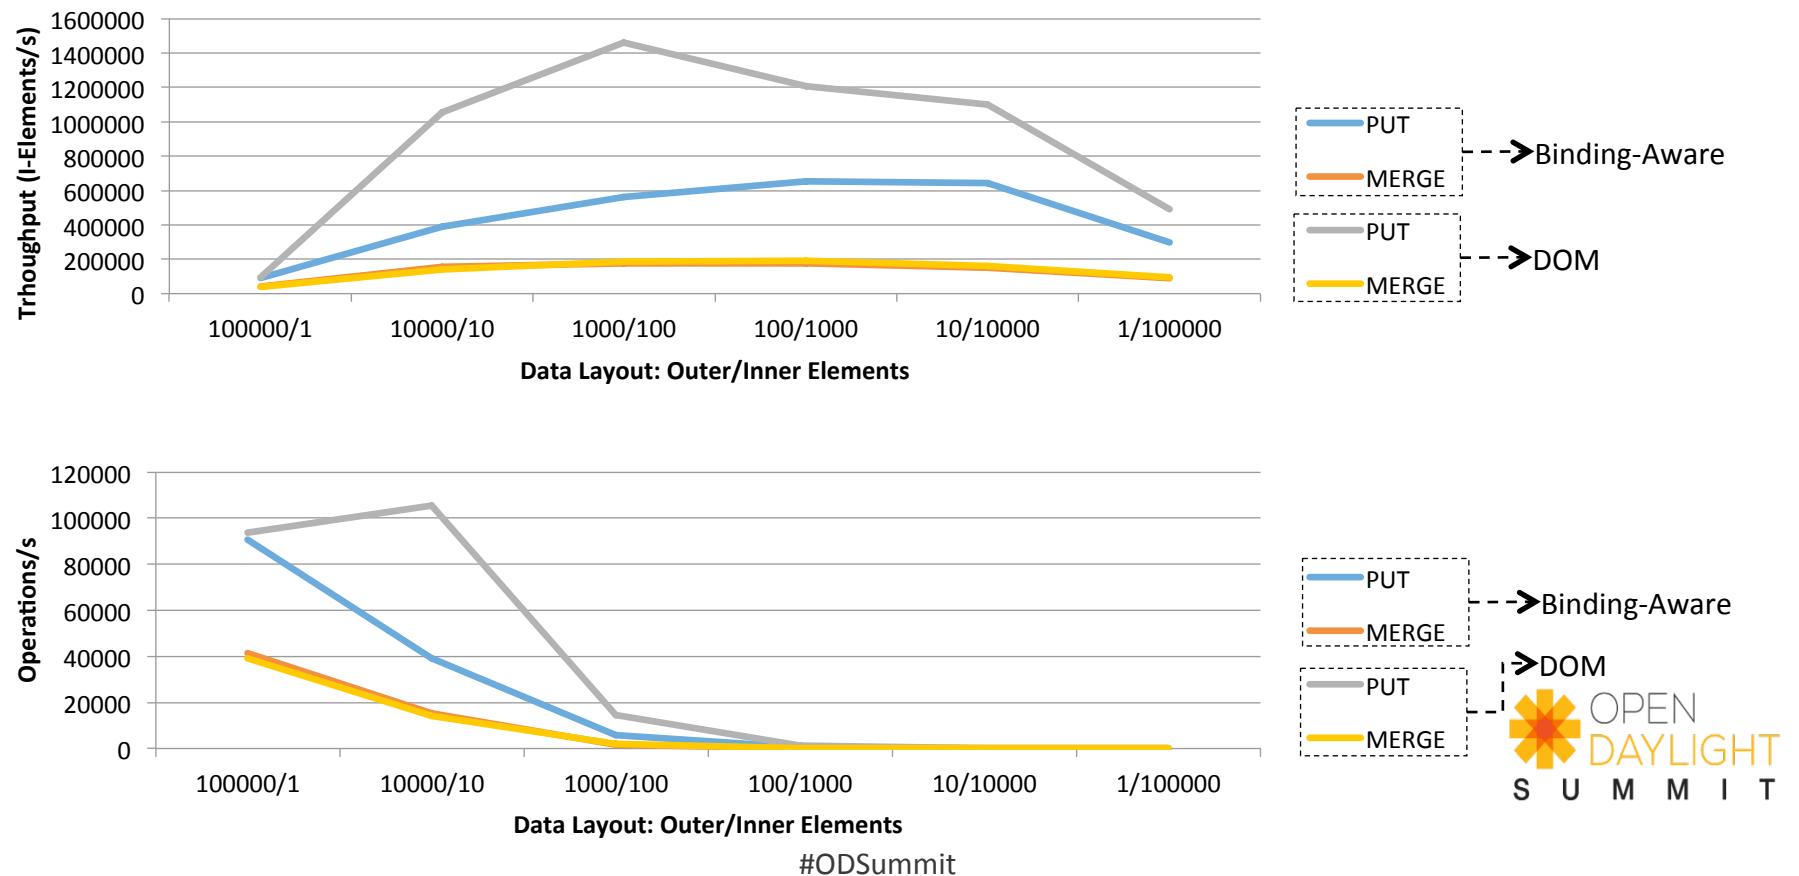
Data Store Benchmark: Persistence

Persistence on/off, transaction chain, pingpong buffer



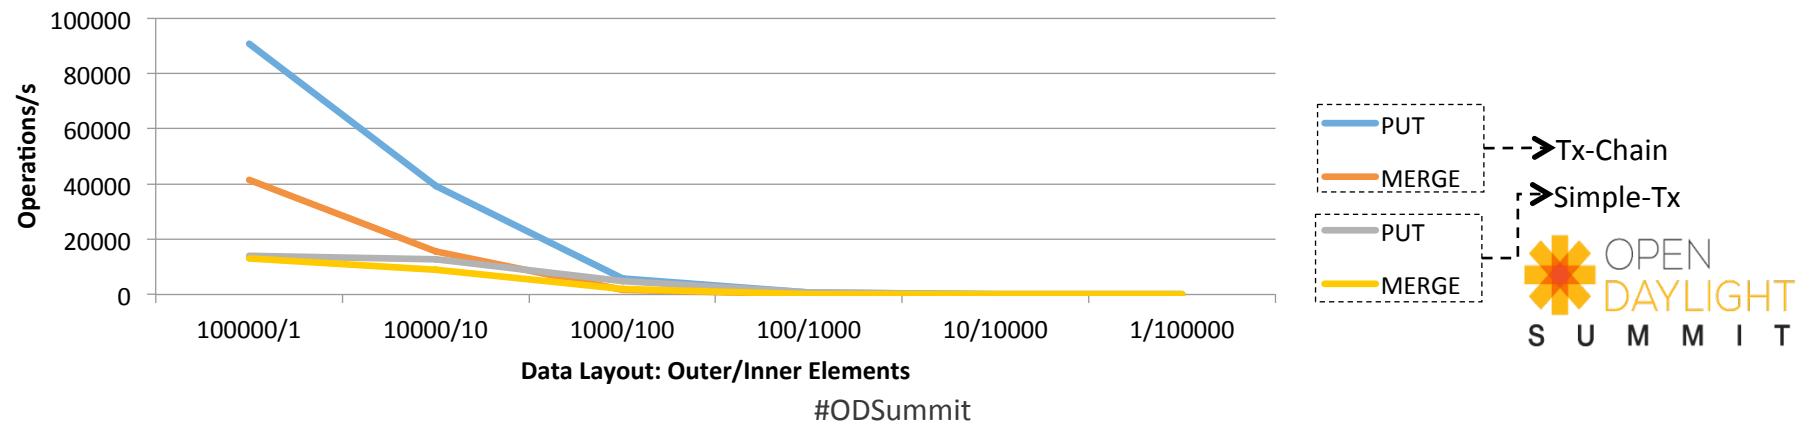
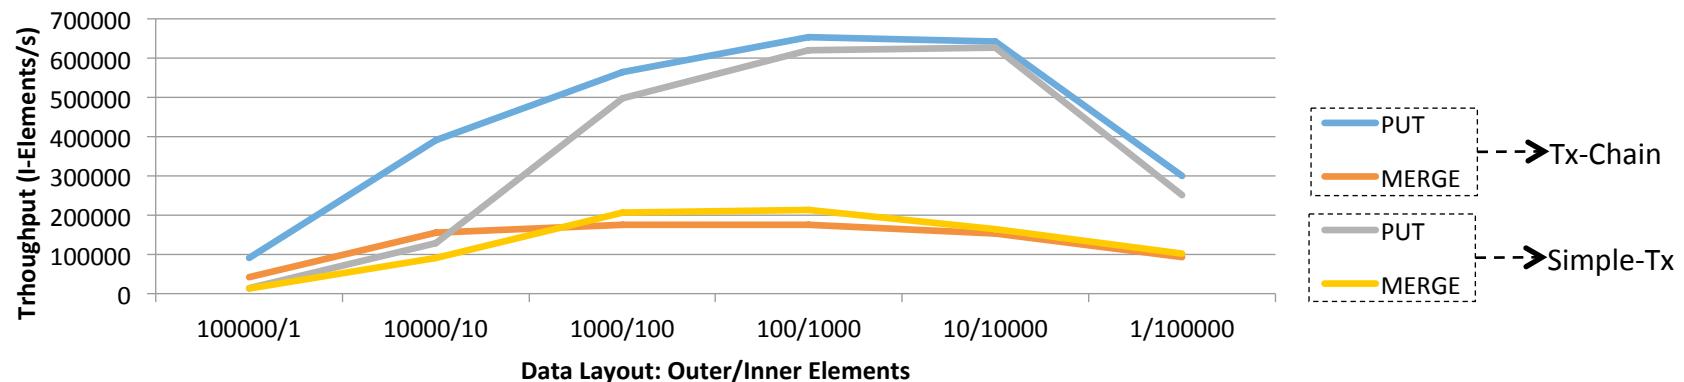
Data Store Benchmark: Binding-Aware vs. DOM

Persistence off, transaction chain, pingpong buffer



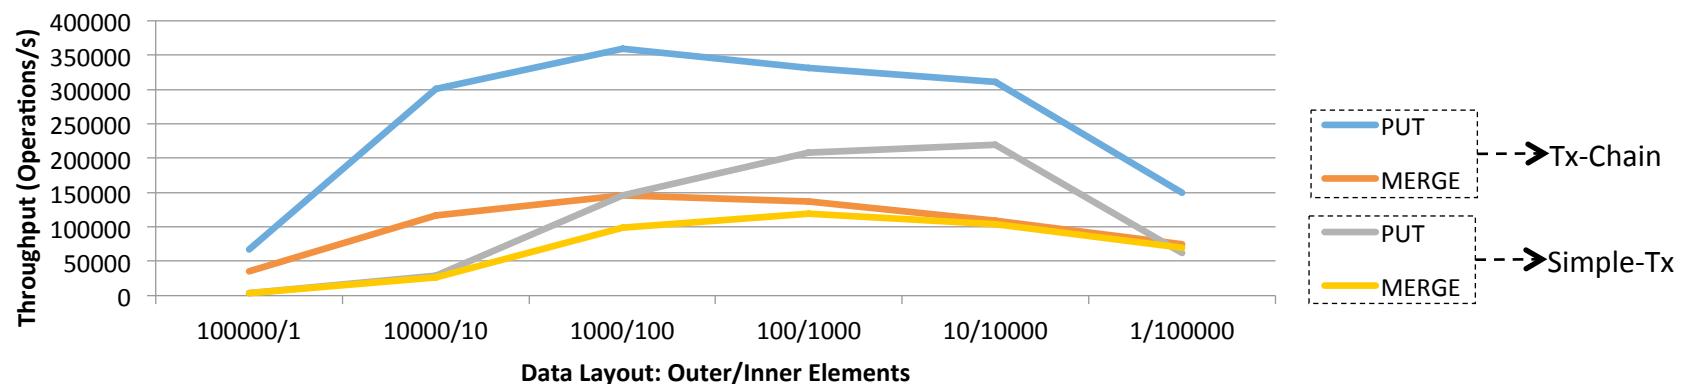
Data Store Benchmark: Simple-Tx vs. Tx-Chain

(**Persistence off**, binding-aware, transaction chain with pingpong buffer)

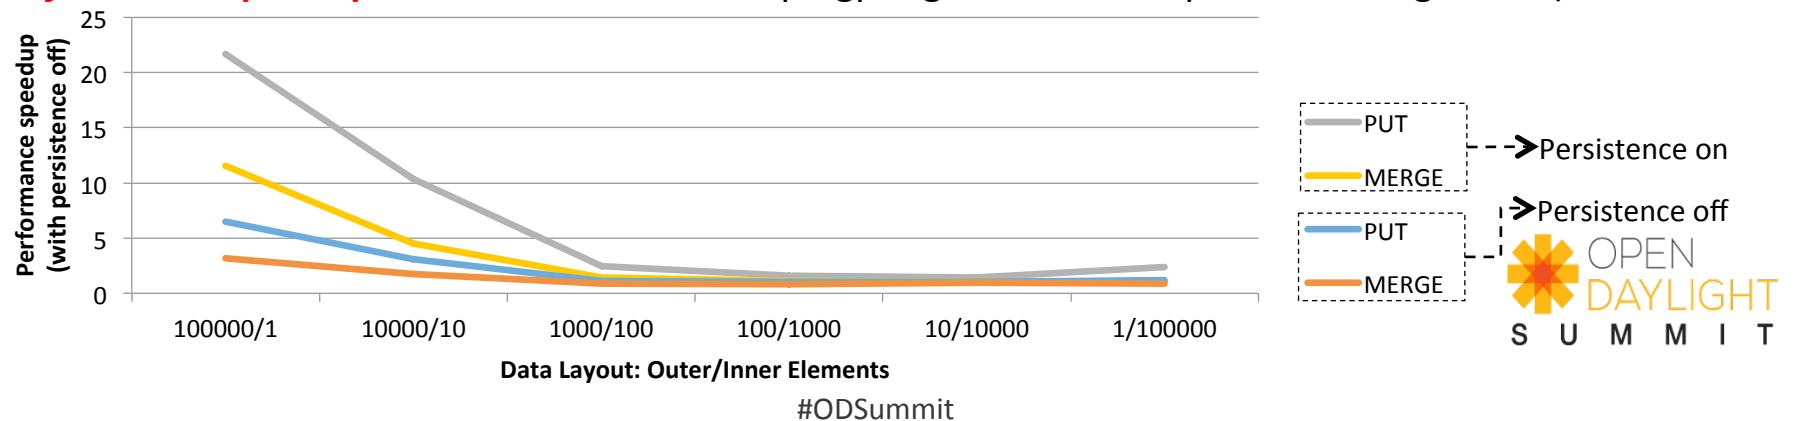


Data Store Benchmark: Simple-Tx vs. Tx-Chain

(*Persistence on*, binding-aware, transaction chain with pingpong buffer vs simple-tx)

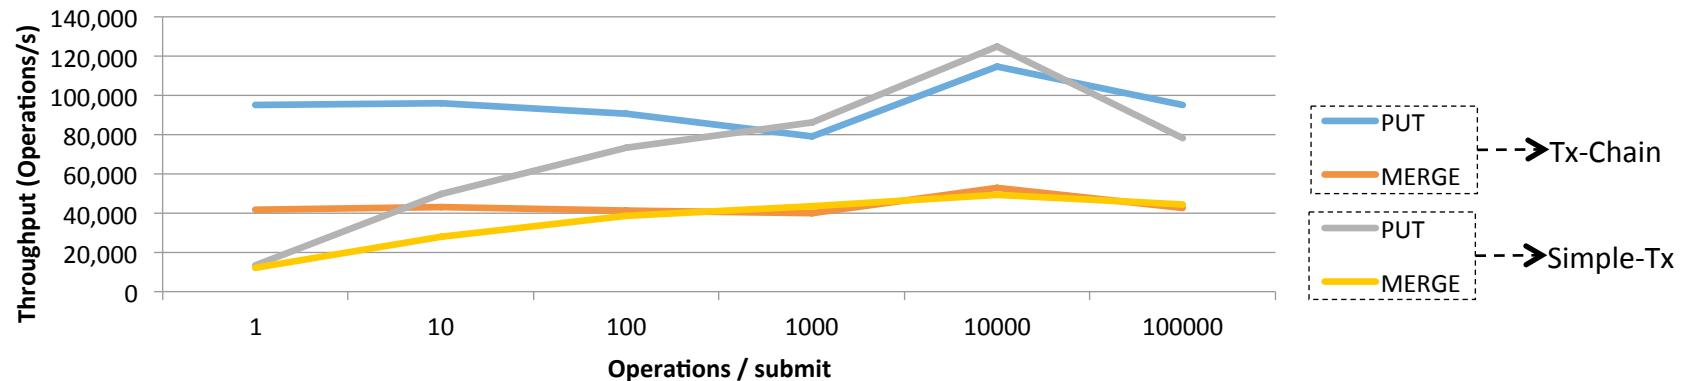


(*Performance speedup* transaction chain with pingpong buffer vs. simple-tx, binding-aware)

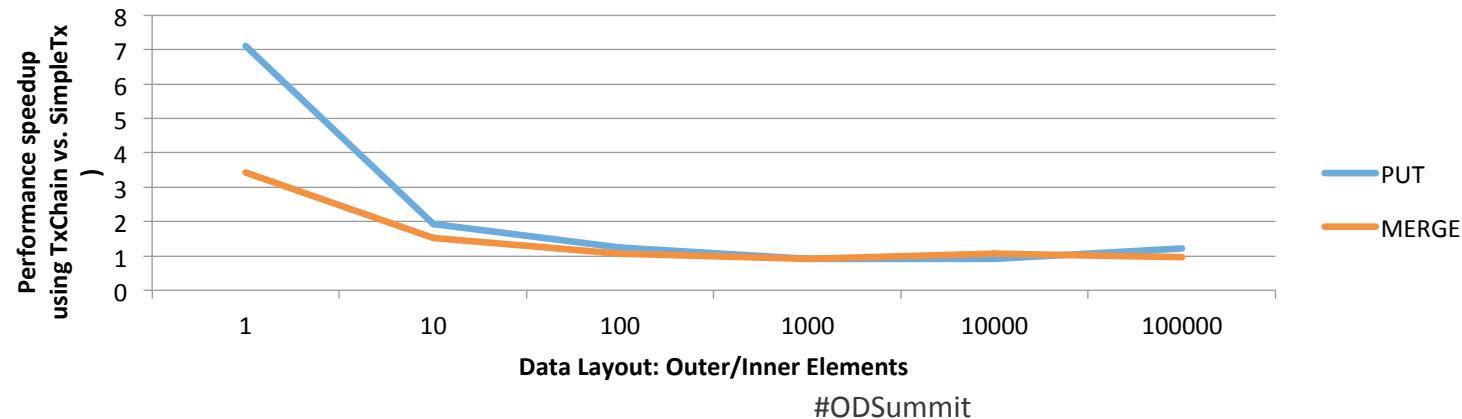


Data Store Benchmark: Batching Operations

100k/1 list, **Persistence off**, Tx-Chain vs. Simple Tx, Binding-Aware

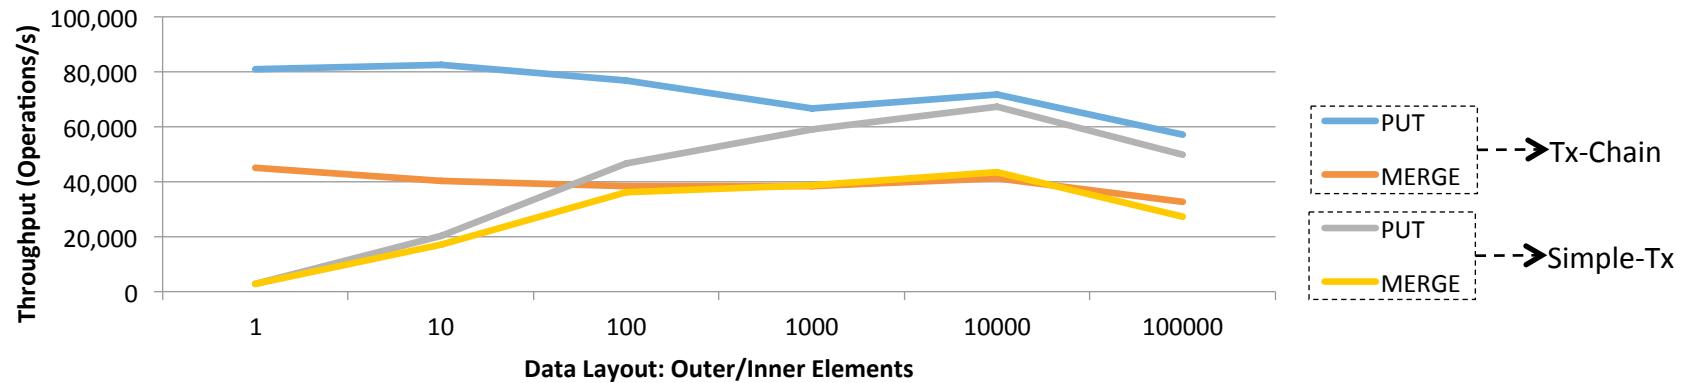


100k/1 list, **Performance speedup**, Persistence off, Tx-Chain vs. Simple-Tx, Binding-Aware

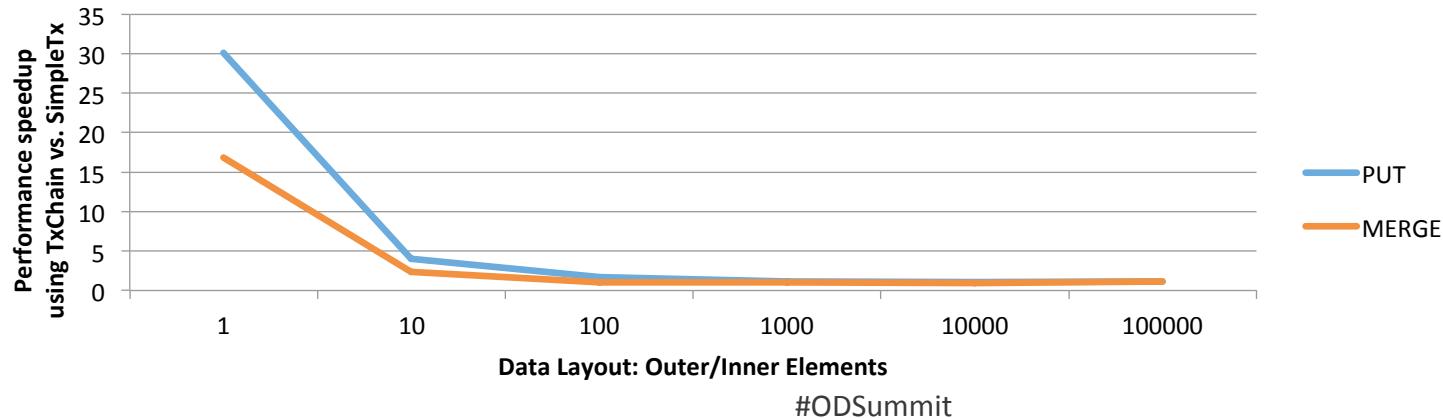


Data Store Benchmark: Batching Operations

100k/1 list, *Persistence on*, Tx-Chain vs. Simple Tx, Binding-Aware



100k/1 list, *Performance speedup*, Persistence on, Tx-Chain vs. Simple-Tx, Binding-Aware





MD-SAL Messaging: RPCs and Notifications

#ODSummit



Using RPCs (Client)

- **Do not block on returned RPC future; use `Futures.addCallback()` instead.**
- Assume that the invocation of an RPC may throw the `RuntimeException` on an incorrect input -> handle if necessary.



#ODSummit

Implementing RPCs (Server)

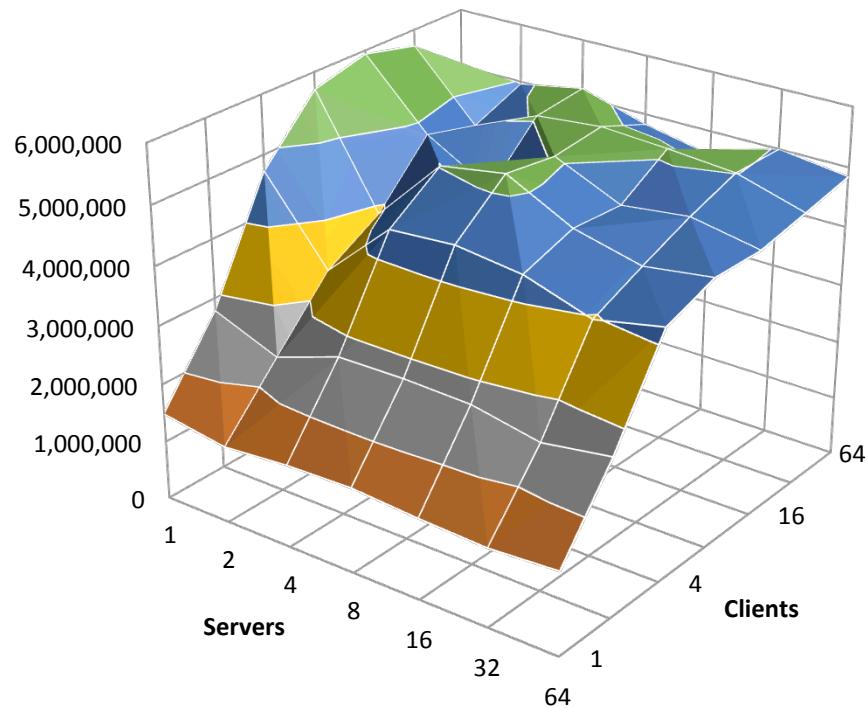
- Do not do CPU-heavy operations in the RPC callback:
 - Dispatch and process requests asynchronously
 - For simplicity use `ListeningExecutorService`
- Where possible return `ListenableFuture`
 - Consider using `SettableFuture` if delegating to a component that has a different asynchronous / callback mechanism.



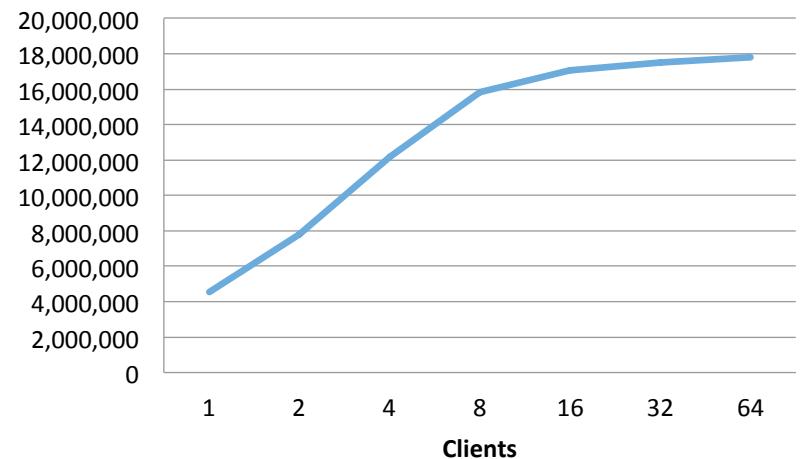
#ODSummit

RPC Benchmark

Routed RPCs



Global RPCs



#ODSummit

Consuming Notifications

- **Do NOT perform any CPU-heavy operations in your callback**
 - Dispatch and process notifications asynchronously if CPU-heavy processing is needed
 - For simplicity use ListeningExecutorService
 - Use advanced dispatch methods if you need more control



#ODSummit

Publishing Notifications

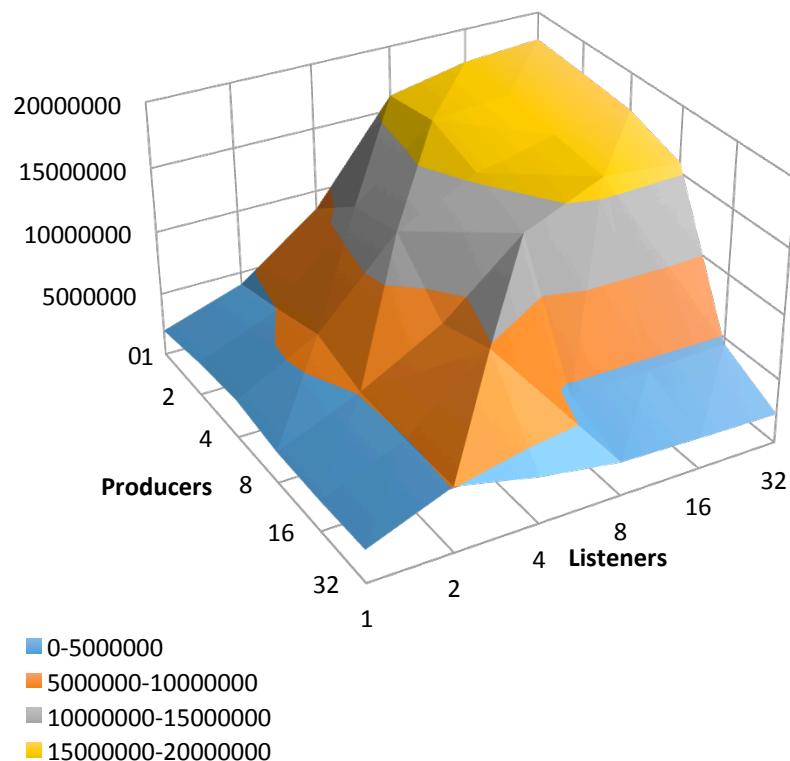
- Use **blocking APIs** (`putNotification()`), only if necessary.
 - **DO NOT** use **blocking APIs** in callbacks from MD-SAL.
- Always use `offerNotification()` for **non-critical notifications** (i.e. those that can be dropped)
 - If you get back the **NotificationPublishService.REJECTED** future, the publishing queue is full -> **back-off** or **slow-down**
 - A future completes once notifications are delivered to the immediate listeners (in the same JVM); may be useful for back-off implementation



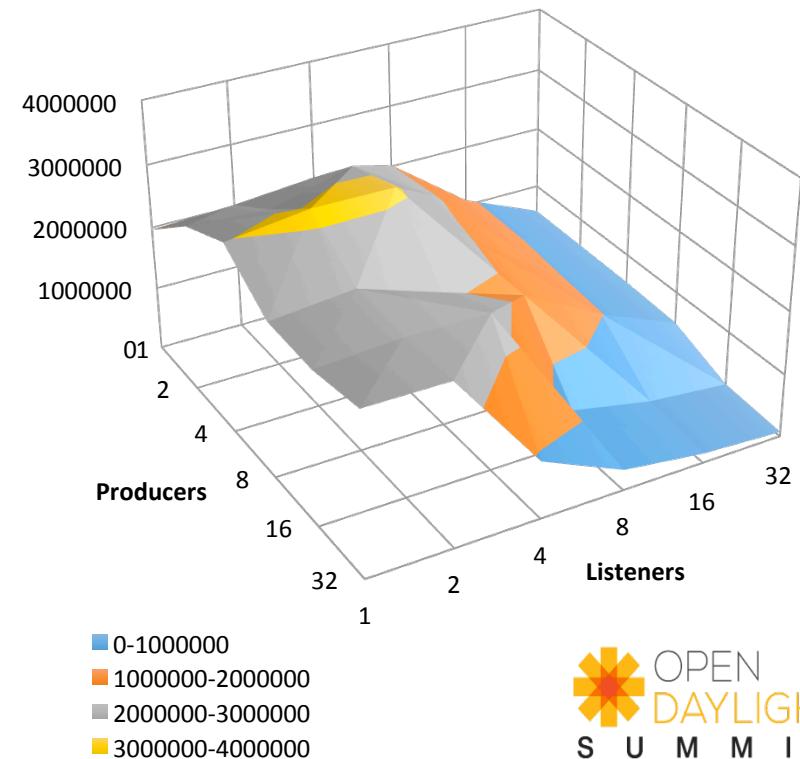
#ODSummit

Notifications Benchmark

Listener Performance



Producer Performance



#ODSummit





Coming in Beryllium

#ODSummit



What? Yet another of new MD-SAL APIs?

- NO, just **improved APIs**, which will allow the MD-SAL to **have a better insight into how data** is produced and consumed:
 - Allows for **better under-the-hood optimization** of data flows
 - Captures the causality of data changes in the system
 - Introduces **micro-sharding**
 - Introduces **more usable APIs**:
 - For **data change listeners** based on feedback from applications
 - For **transaction chaining**
 - For **operation batching**
 - **Actor-like model** for listeners / producers



#ODSummit

How Will This Affect My Code?

- Affects the **instantiation** of listeners and transactions, not their use:
 - The new APIs have a lot in common with the current APIs
 - The concepts are similar
 - Migration Guide on the way
- **Examples:**
 - **DataTreeChangeListener -> DataChangeListener**
 - You get all the state that you registered for – **no need for READs**
 - **TransactionChain -> DataTreeProducer**
 - Subtree scoped, easier to use, allows for opt-in / opt-out of batching of transaction operations



#ODSummit



Questions?

#ODSummit



Where to Find More Information

- Details:
[https://wiki.opendaylight.org/view/
Controller_Core_Functionality_Tutorials:Tutorials:Data_Store_Benchmarking_and_Data_Access_Patterns](https://wiki.opendaylight.org/view/Controller_Core_Functionality_Tutorials:Tutorials:Data_Store_Benchmarking_and_Data_Access_Patterns)
- Code: coretutorials/dsbenchmark



#ODSummit



Thank You



#ODSummit



Test Data: List of Lists

Yang Model:

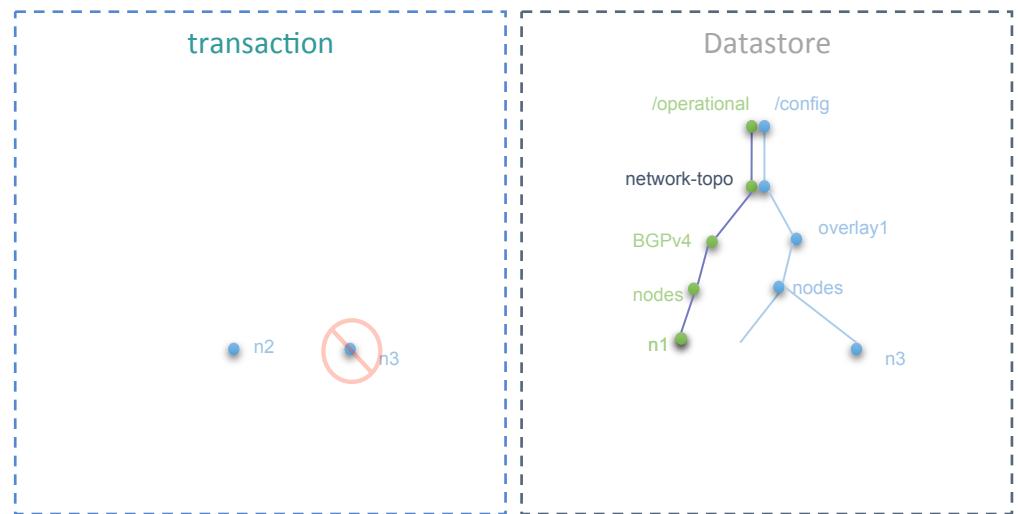
```
container test-exec {  
    config true;  
  
    list outer-list {  
        key id;  
        leaf id {  
            type int32;  
        }  
  
        list inner-list {  
            key name;  
            leaf name {  
                type int32;  
            }  
  
            leaf value {  
                type string;  
            }  
        }  
    }  
}
```

Using Generated Java DTOs:

```
static public List<OuterList> buildOuterList(int outerElements,  
                                              int innerElements) {  
    List<OuterList> outerList = new ArrayList<OuterList>(outerElements);  
    for (int j = 0; j < outerElements; j++) {  
        outerList.add(new OuterListBuilder().setId( j )  
                     .setInnerList(buildInnerList(j, innerElements))  
                     .setKey(new OuterListKey( j )).build());  
    }  
    return outerList;  
}  
  
static private List<InnerList> buildInnerList( int index, int elements ) {  
    List<InnerList> innerList = new ArrayList<InnerList>( elements );  
  
    final String itemStr = "Item-" + String.valueOf(index) + "-";  
    for( int i = 0; i < elements; i++ ) {  
        innerList.add(new InnerListBuilder().setKey( new InnerListKey( i ) )  
                     .setName(i).setValue( itemStr + String.valueOf( i ) )  
                     .build());  
    }  
    return innerList;  
}
```

Transactions: Read/Write/Delete

```
ReadWriteTransaction transaction =  
dataBroker.newReadWriteTransaction();  
Optional<Node> nodeOptional;  
nodeOptional = transaction.read(  
    LogicalDataStore.OPERATIONAL,  
    n1InstanceIdentifier);  
transaction.put(  
    LogicalDataStore.CONFIGURATION,  
    n2InstanceIdentifier,  
    topologyNodeBuilder.build());  
transaction.delete(  
    LogicalDataStore.CONFIGURATION,  
    n3InstanceIdentifier);  
CheckedFuture future;  
future = transaction.submit();
```



Transactions: Merge vs. Put

```
WriteOnlyTransaction transaction =  
dataBroker.newWriteOnlyTransaction();  
InstanceIdentifier<Node> path =  
InstanceIdentifier  
.create(NetworkTopology.class)  
.child(Topology.class,  
new TopologyKey(  
"overlay1"));  
transaction.merge(  
LogicalDataStore.CONFIG,  
path,  
topologyBuilder.build());  
CheckedFuture future;  
future = transaction.submit();
```

