

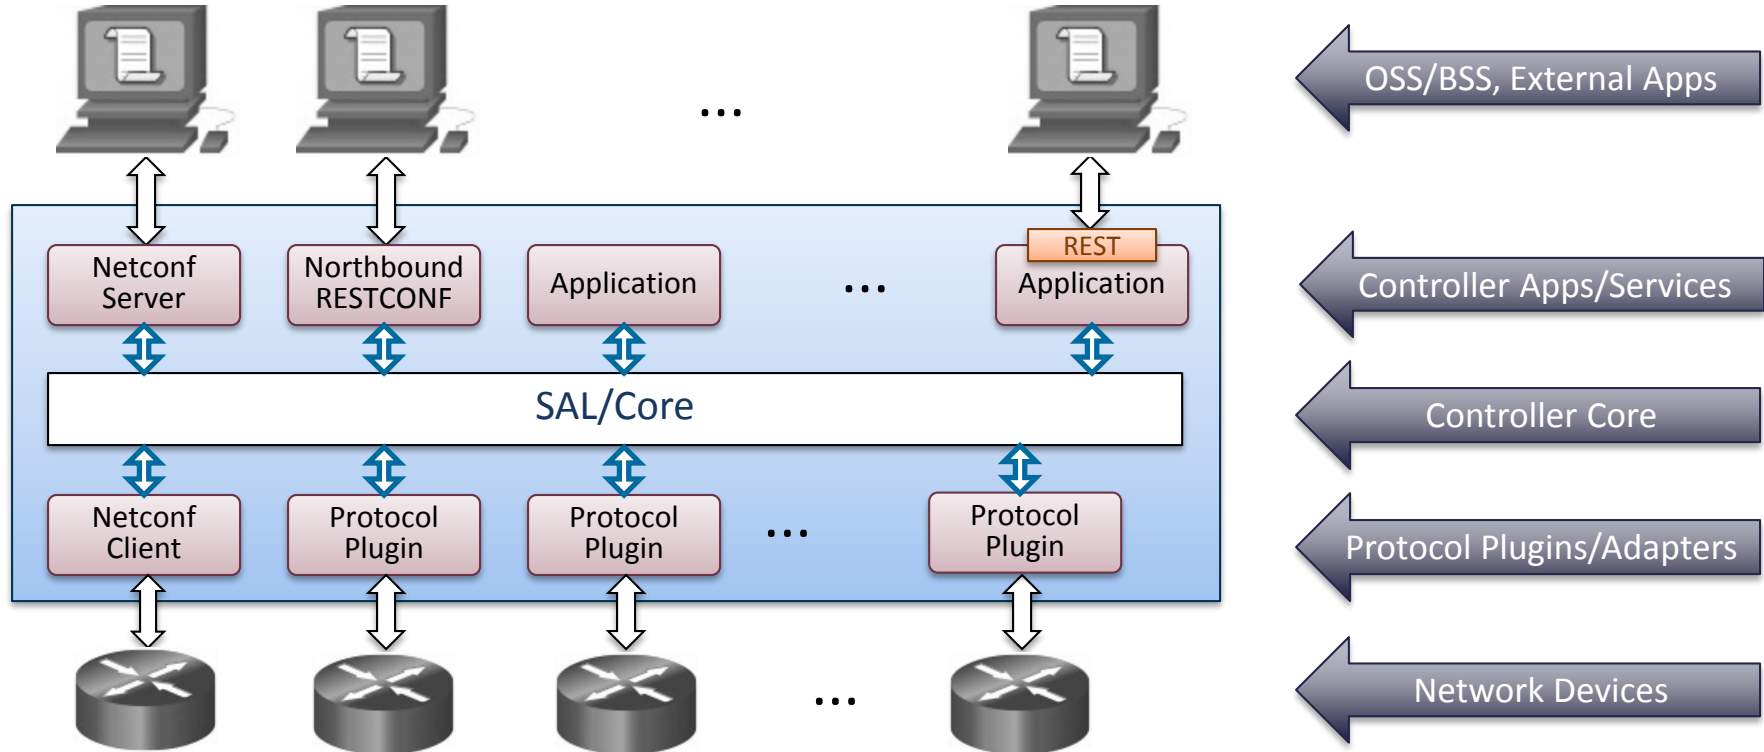
OpenDaylight YANG Data Store: A High-Performance Data Store for SDN and IoT Applications

Jan Medved

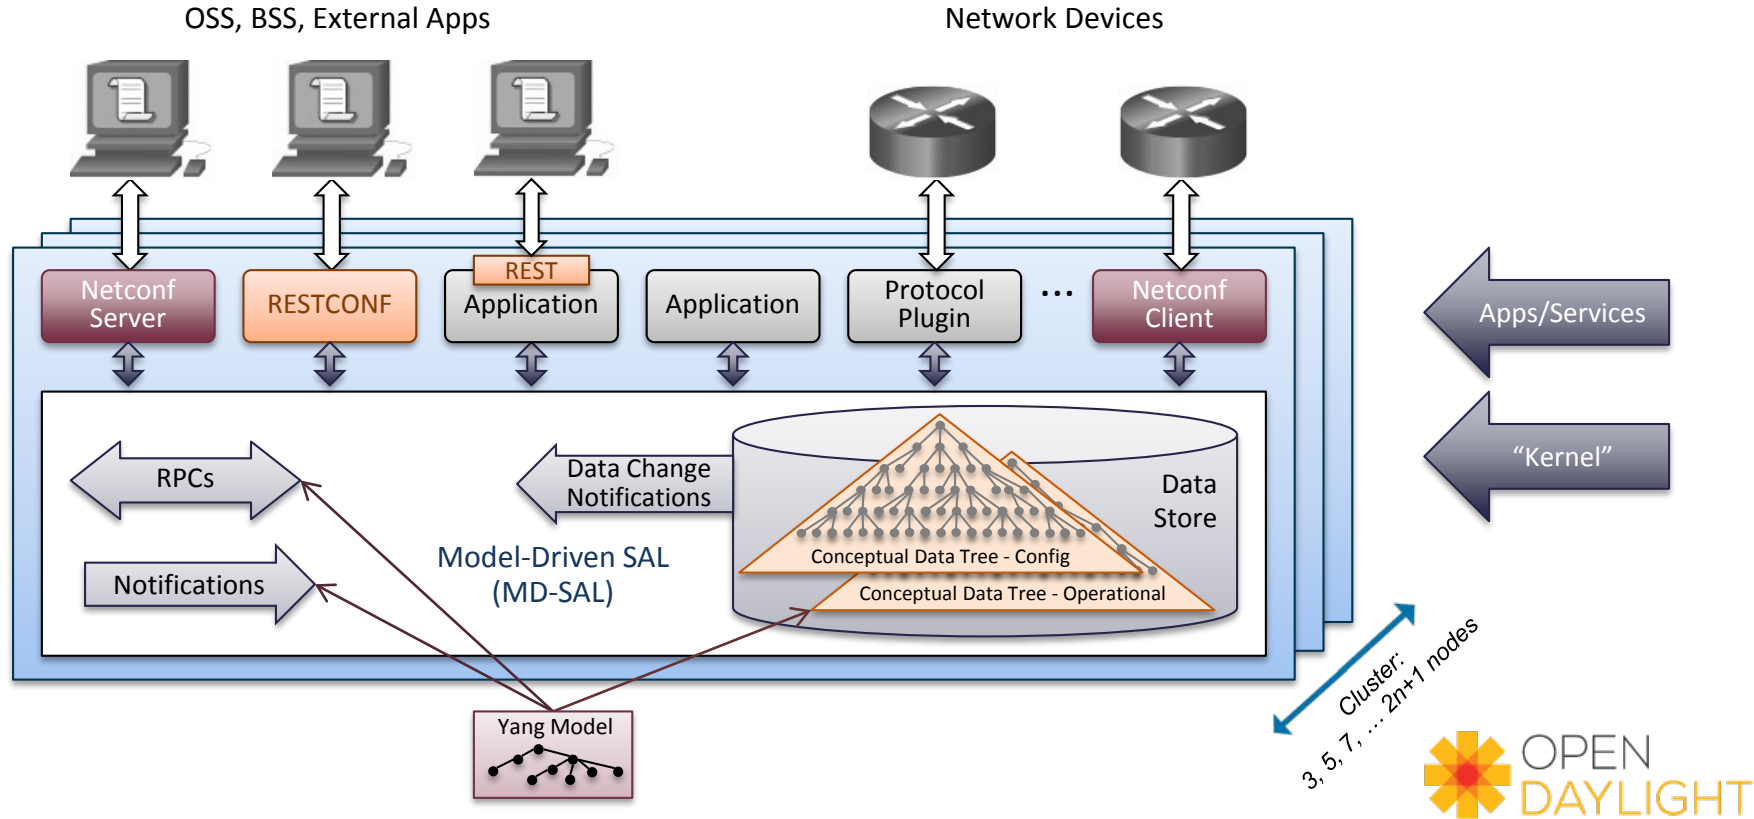
With contributions from Robert Varga, John Burns, Kun Chen, Branislav Janosik



OpenDaylight is an SDN Controller, Right?



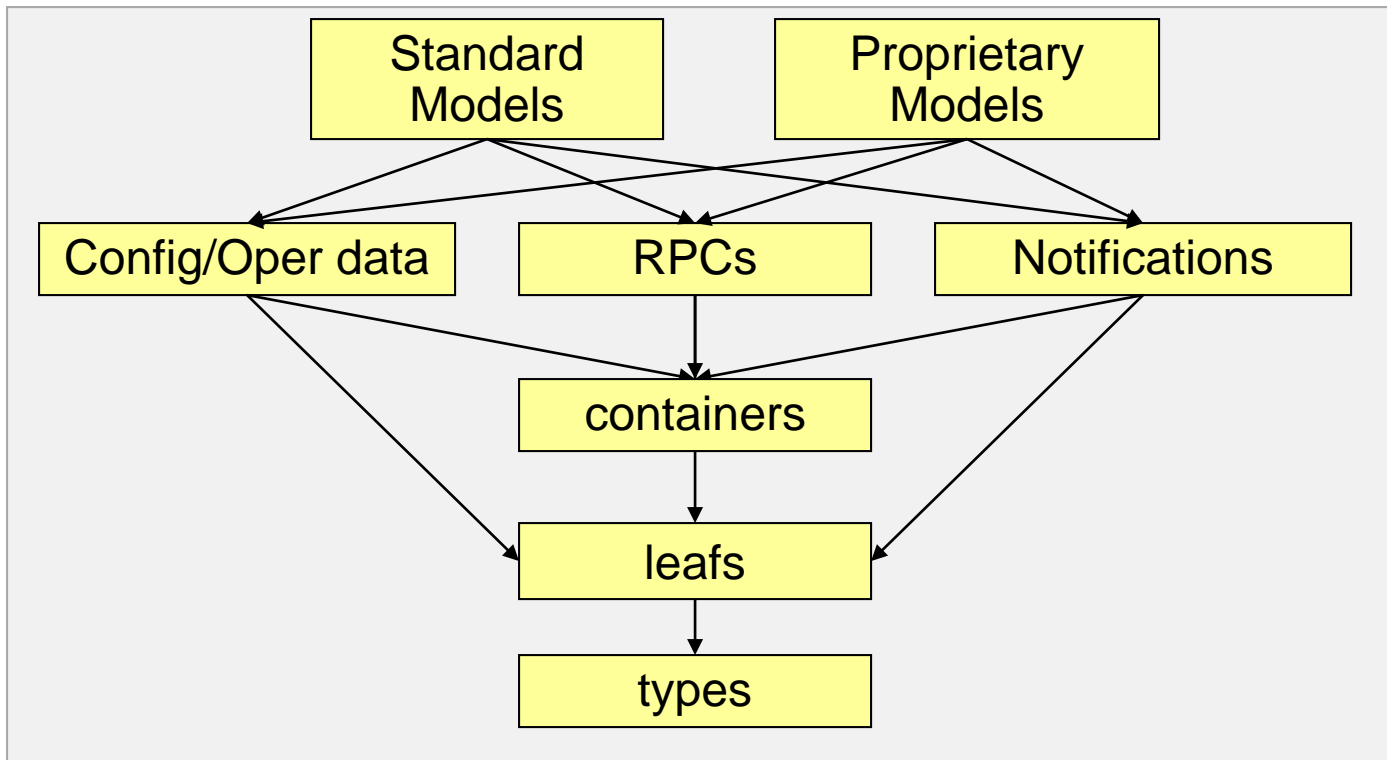
ODL Architecture: an Application Development Platform



YANG is

- A NETCONF modeling language
 - Think SMI for NETCONF
- Models semantics and data organization
 - Syntax falls out of semantics
- Able to model config data, state data, RPCs, and notifications
- Based on SMIng syntax
 - Text-based
 - Email, patch, and RFC friendly
- *Also used as an Interface Description Language in OpenDaylight*

YANG Concepts




YANG

- Directly maps to XML or JSON content on the wire
- Extensible
 - Add new content to existing data models
 - Without changing the original model
 - Add new statements to the YANG language
 - Vendor extensions and future proofing
- Tools in OpenDaylight to generate Java Code from yang models
 - Compile and runtime
- See tools at www.yang-central.org

Yang Data Store Challenge: Conceptual Data Tree

Yang defines a tree address space:

- Single inheritance for containers and leafs






The OpenDaylight Data Store...

- Implements the equivalent of a W3C DOM Document tree
 - No parent axis maintained -> efficient copy-on-write snapshot
- MVCC-based atomic updates in a single-writer, multiple-readers environment
 1. Modifications prepared concurrently
 2. Single update thread:
 - Requested mods are applied and made visible to subsequent snapshots
 3. Efficient physical replication
- Enforcement of structural integrity according to yang schemas


Data Store Operation:

Step 1- Transaction Created

Initial snapshot when transaction created



Modification tree (populated as mods in the transaction are performed)




Data Store Operation:

Step 2 - Creating a New Node (Write)

Initial snapshot when transaction created

Modification tree (populated as mods in the transaction are performed)




Data Store Operation:

Step 3 - Transaction Seal (Ready)

Initial snapshot when transaction created

Modification tree (populated as mods in the transaction are performed)




Steps 2 and 3 can run concurrently in multiple threads

Data Store Operation:

Step 4a - Transaction Prepare (Walk Down)

Initial snapshot when transaction created

Modification tree (populated as mods in the transaction are performed)



Data Store Operation:


Step 4b -Transaction Prepare (Walk Up)

Initial snapshot when transaction created

Modification tree (populated as mods in the transaction are performed)


Prepare transaction:

- Create a copy of the current tree (only copy modified nodes)
- Retain old data in nodes (both new and old data available in current state)
- Exception if data does not conform to schema



Creating this node in the tree


Original node pointers



*Logical operation = Touch
Modification Type = Subtree Modified*

*Logical operation = Write
Modification Type = Write*

Steps 4 and 5 must be executed in one and the same thread



Data Tree Candidate

Current state of the tree

Data Store Operation:


Step 5: Commit

Steps 4 and 5 must be executed in one and the same thread

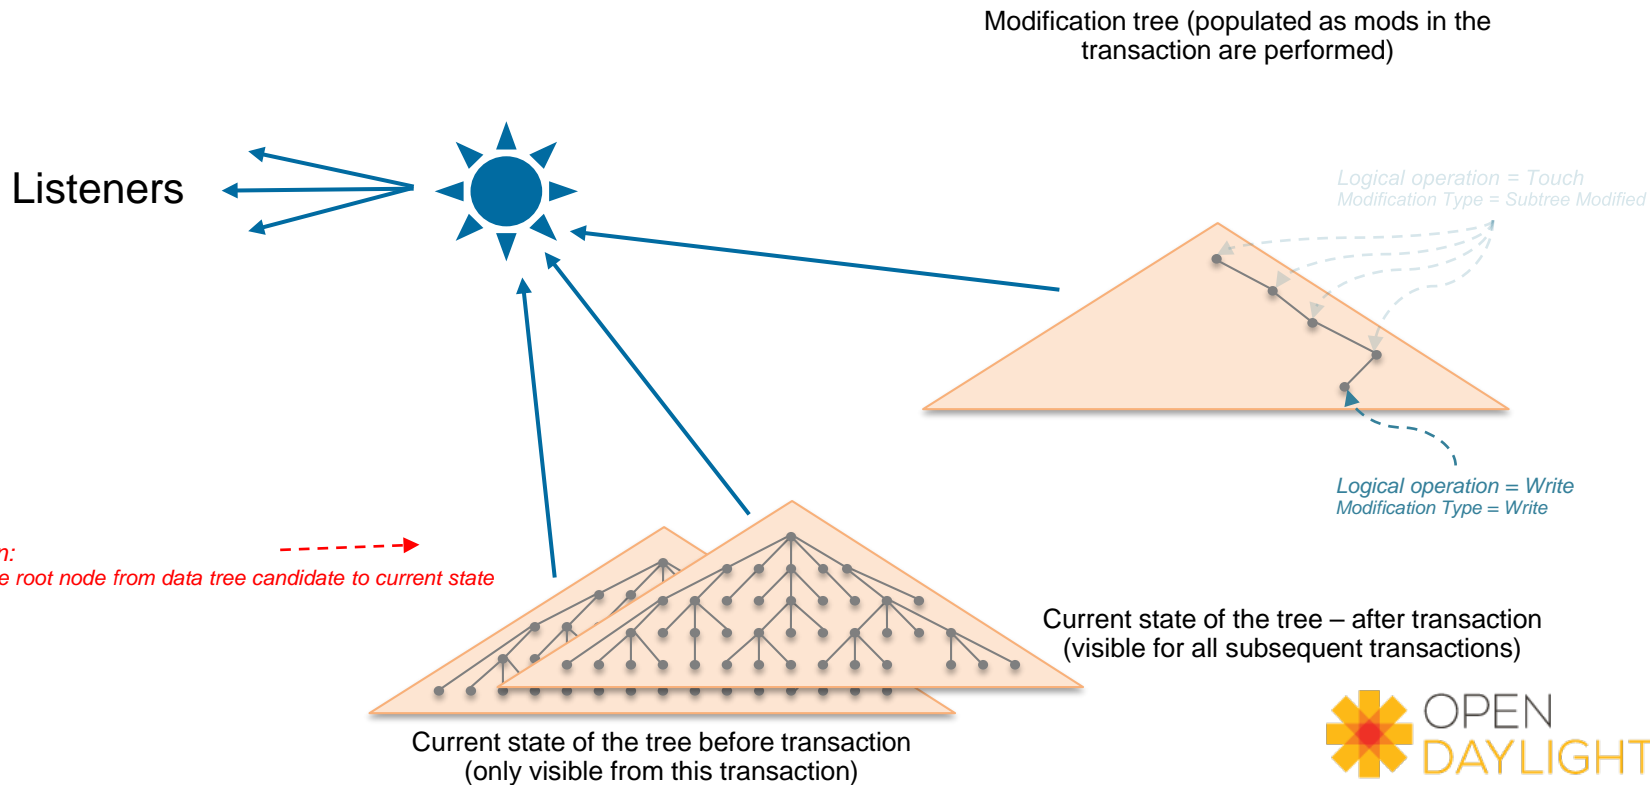
Initial snapshot when transaction created:

- If commit successful, only used in transactions that were created on top of the just committed transaction
- will be garbage collected if not needed anymore

Modification tree (populated as mods in the transaction are performed)



Step 6: Notify Listeners

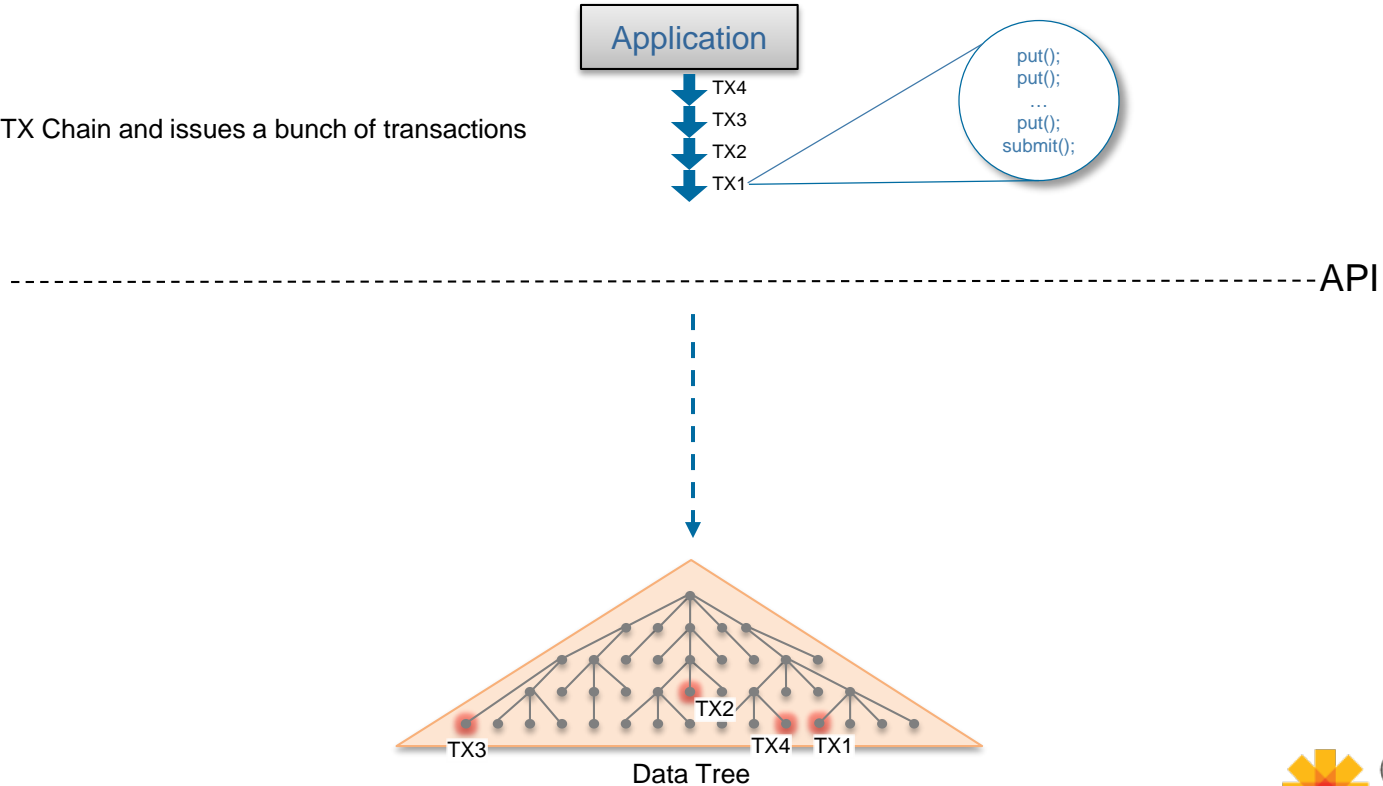


How to Improve Transaction Performance

- State Compression
 - Multiple fine-grained transactions compressed into one bigger transaction
 - “Fate sharing” -> faster success path, but longer recovery
 - Transaction Chaining with ping-pong buffer
 - Single writer
- Sharding (Parallelism)

Transaction Chaining with Ping-Pong Buffer

1. App creates a TX Chain and issues a bunch of transactions



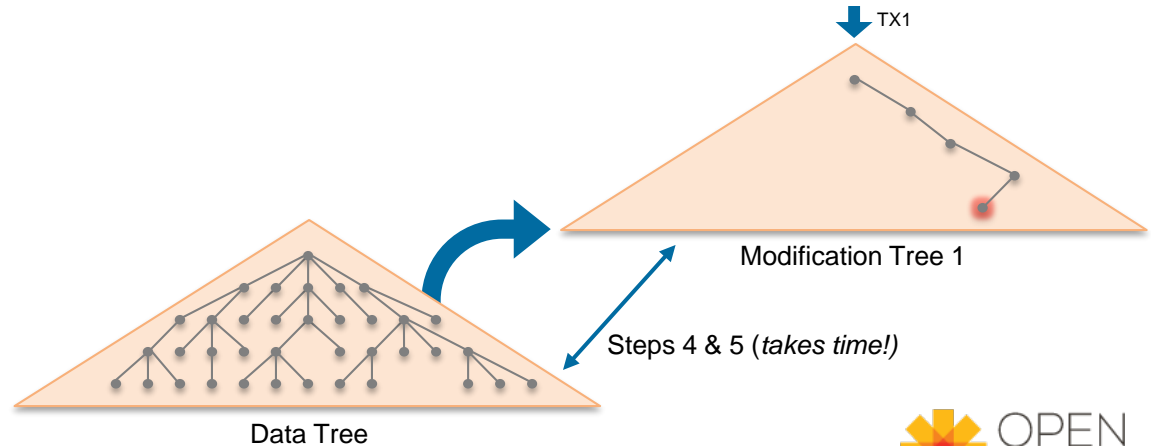
Transaction Chaining with Ping-Pong Buffer

Application

↓ TX4
↓ TX3
↓ TX2

2. First transaction creates a Modification Tree
- submit() starts Steps 4-5

-----API




Transaction Chaining with Ping-Pong Buffer

3. TX2 issued while TX1 submit() being processed

Application

TX4
TX3

API

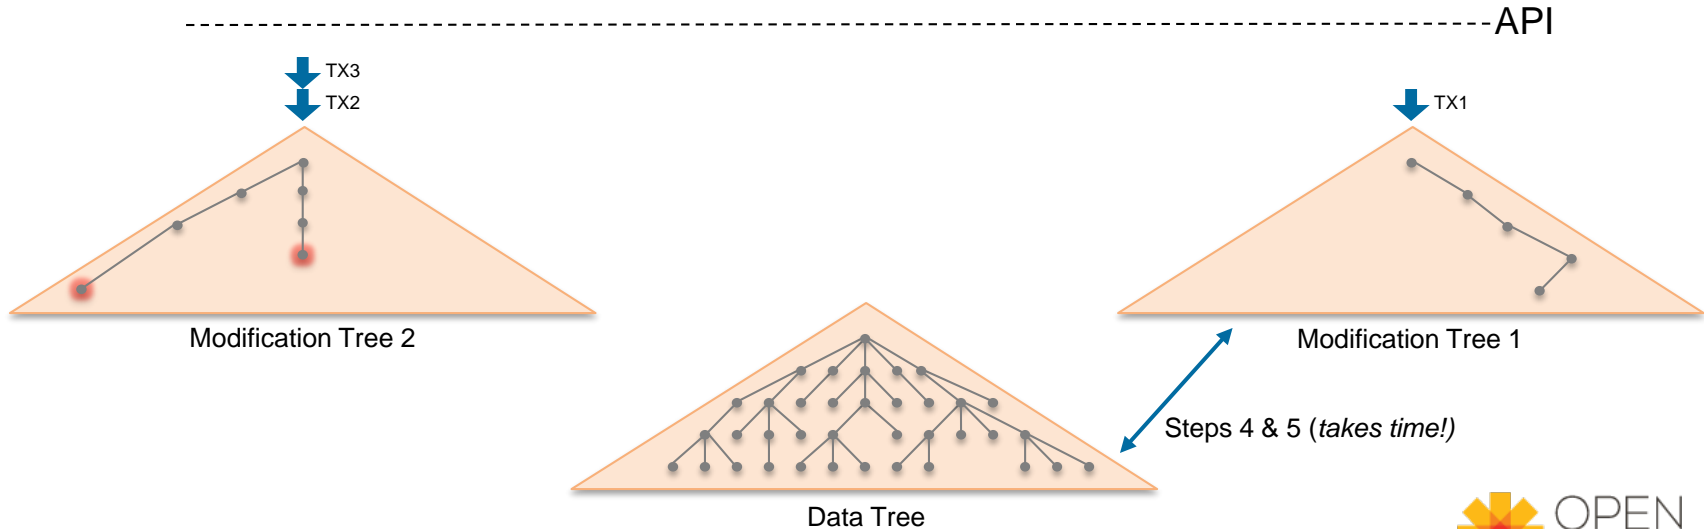


Transaction Chaining with Ping-Pong Buffer

3. TX3 issued while TX1 submit() is STILL being processed

Application

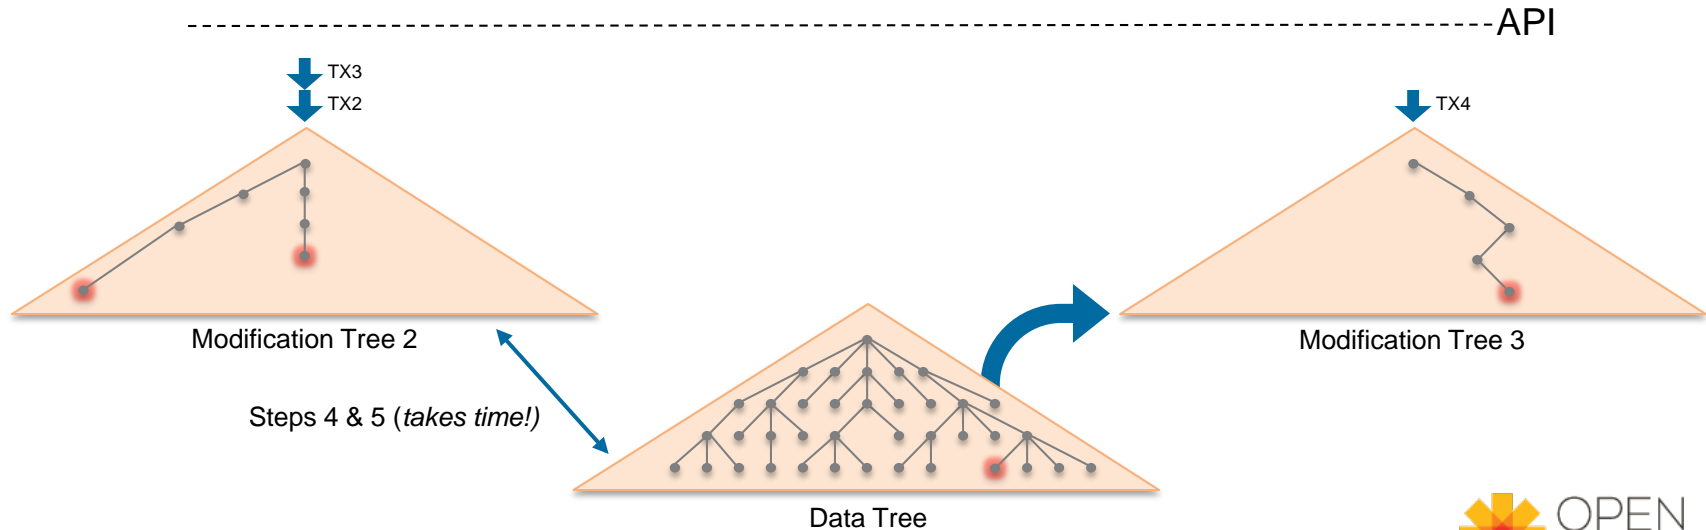
↓ TX4




Transaction Chaining with Ping-Pong Buffer

Application

4. TX4 issued while combined TX2 & TX3 being submitted

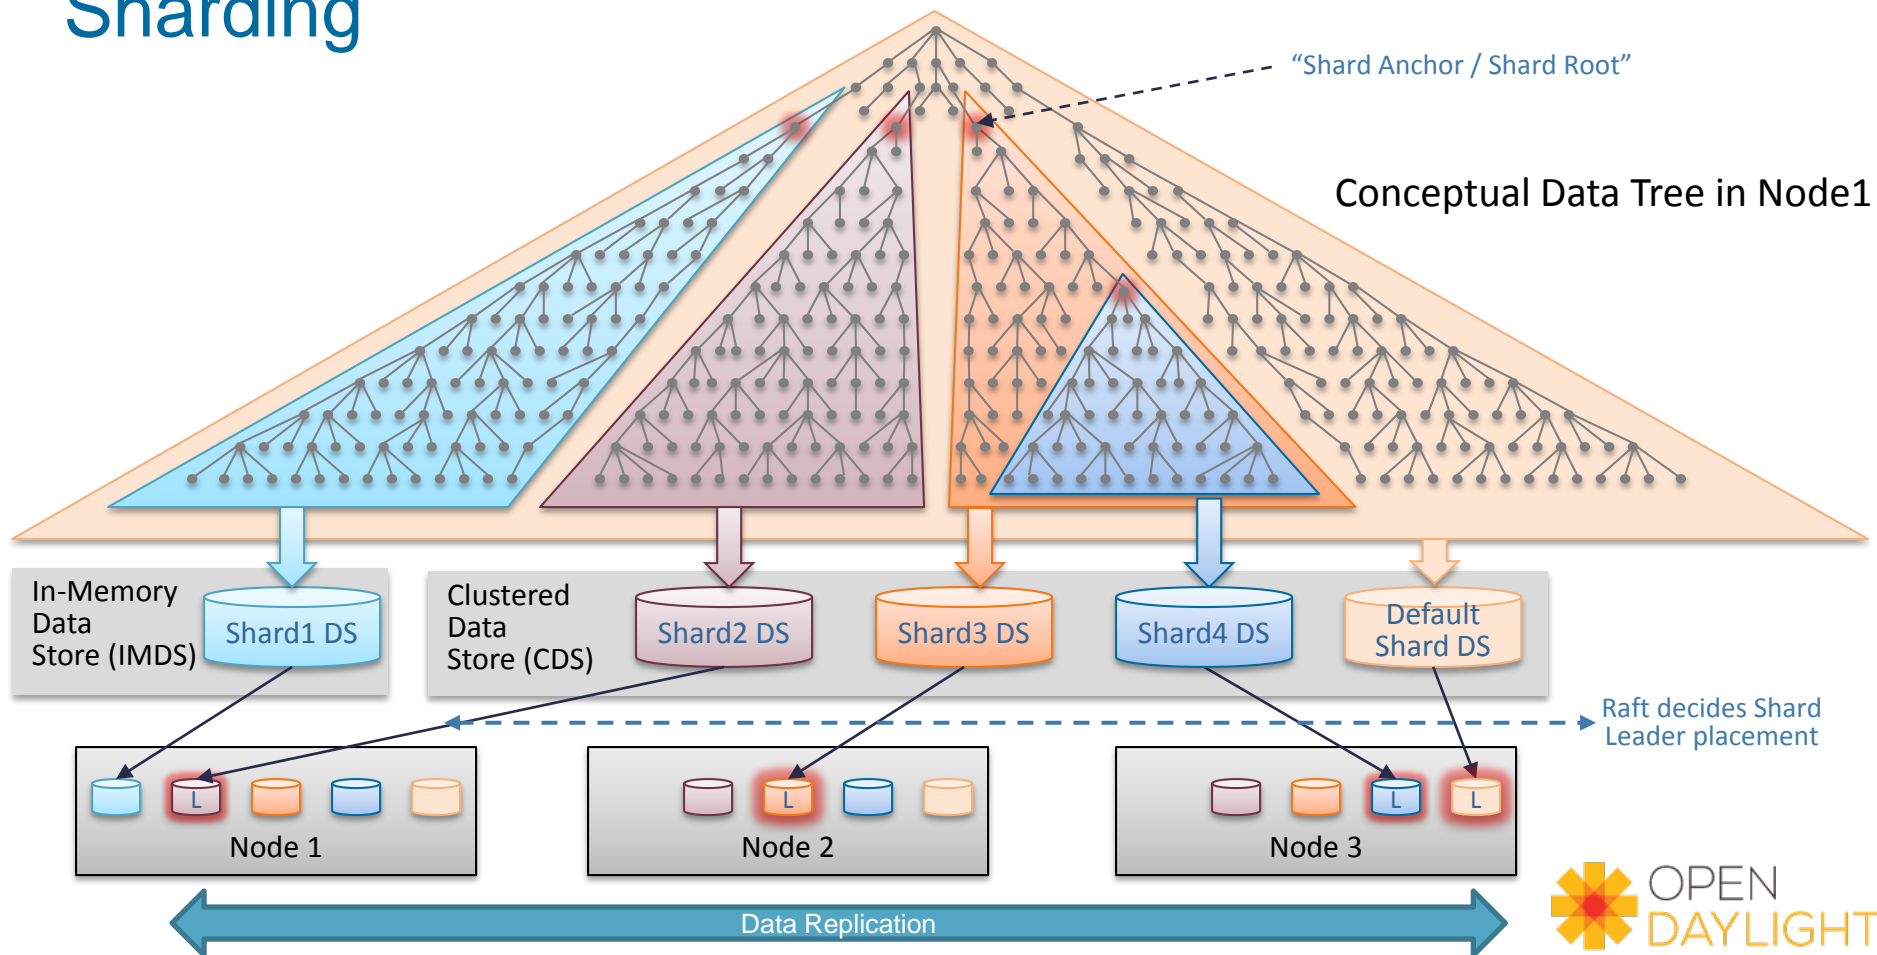


Transaction Chain vs Single Transactions Performance

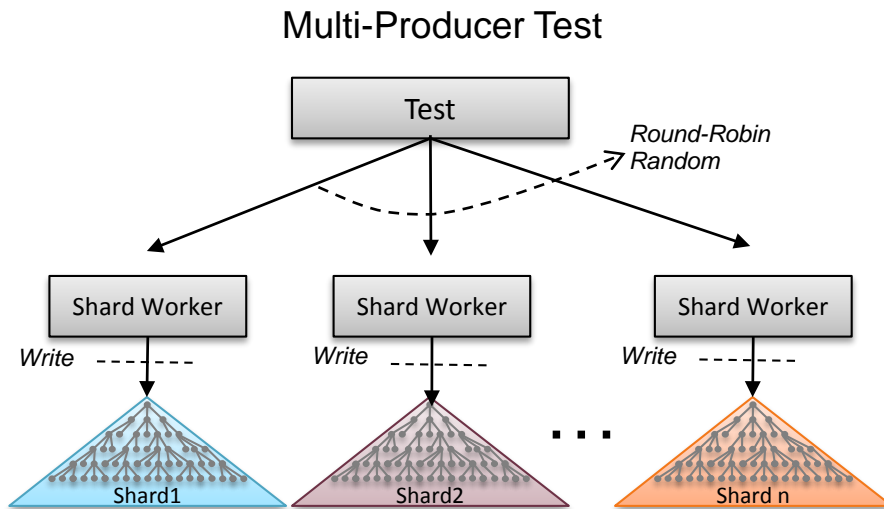


Test Data: List of lists

Sharding




Shard Performance Testing: Test Setup



- List/Array of x string elements divided into 1 ... n shards
- (Test) writes data to Shard Workers (Producers) in Round-Robin or Random fashion
- Worker handles tx responses asynchronously (Futures)

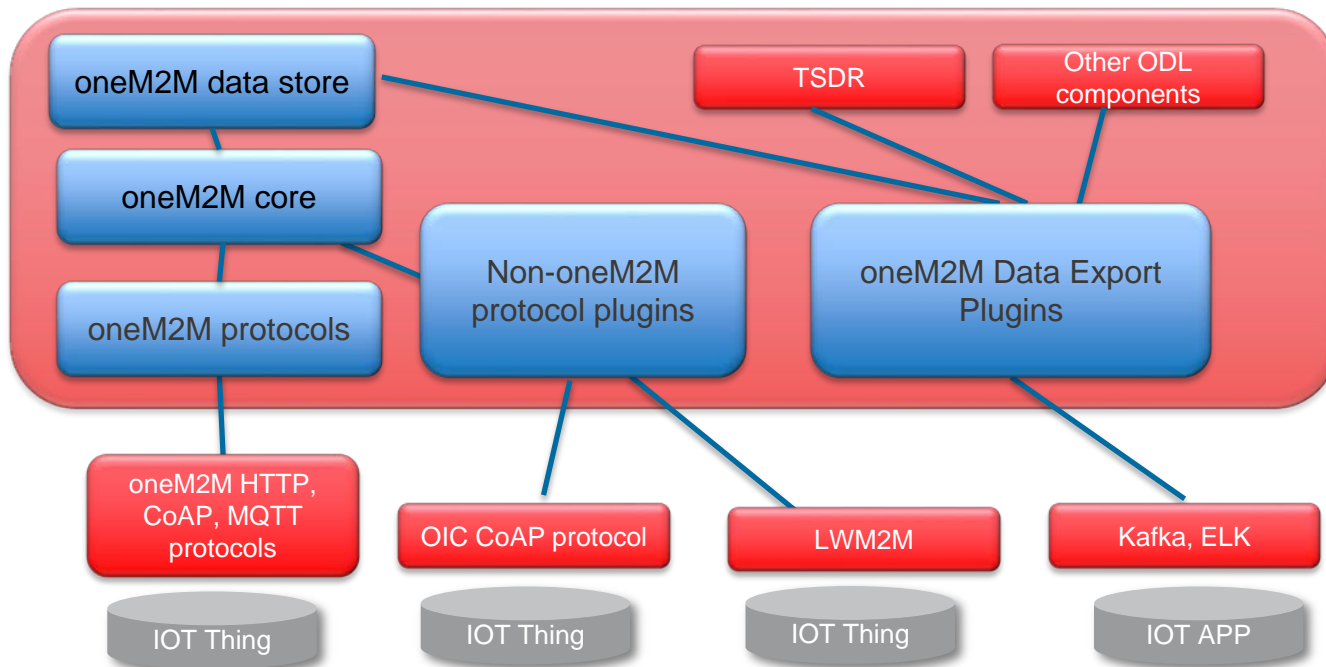
Performance Results: Multiple Producers




IoTDM

- What is IoTDM?
 - A collection of projects for IOT data management.
 - git clone <https://git.opendaylight.org/gerrit/p/iotdm.git>
 - <https://wiki.opendaylight.org/view/IoTDM:Main>
 - Initial IOT standard chosen to implement is onem2m.org
 - Model IOT things and their data in the ODL data store
 - onem2m resources and resource tree modeled generically in the onem2m.yang file
 - Other standards in the future such as
 - OIC, LWM2M, ...
 - Initial strategy for other standards is to adapt /interwork the other standards, normalized to onem2m

IoTDM Architecture




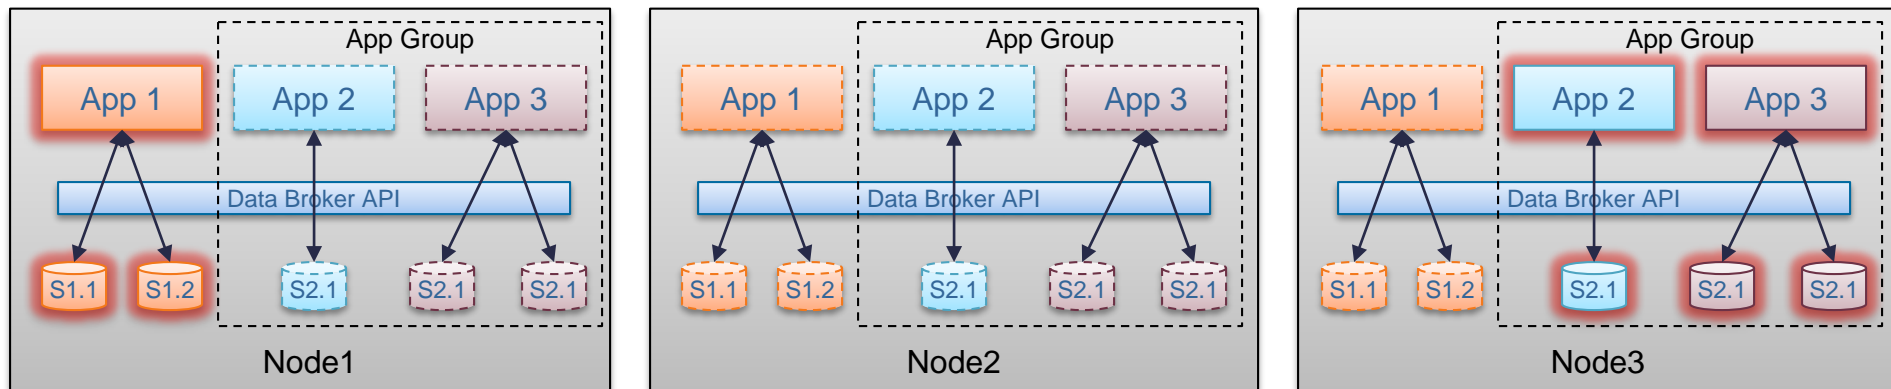
oneM2M Tree for Performance



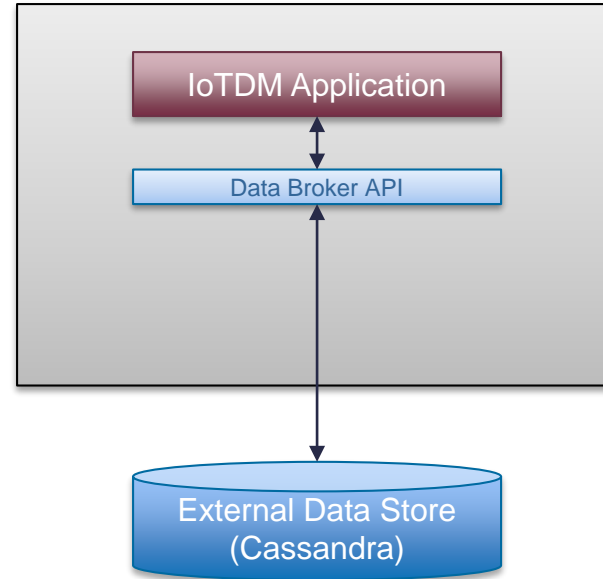
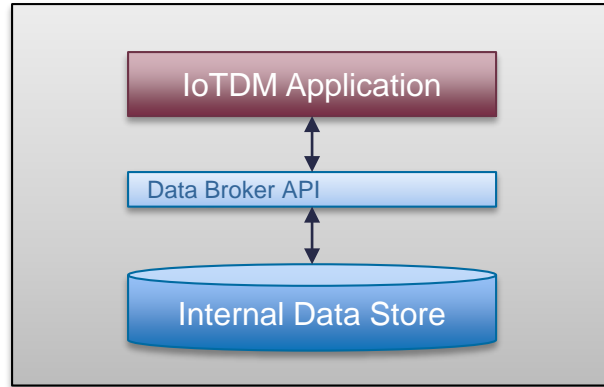
OpenDaylight Application Model

Service/Application:

- Code + Data (Shards, subtrees of the Conceptual Data Tree)
- Service/Application instances SHOULD be co-located with shard replicas
- Active Service/Application instance SHOULD be co-located with all Shard Leaders it “owns”
- Apps can be grouped for “fate sharing




Internal vs. External Data Store

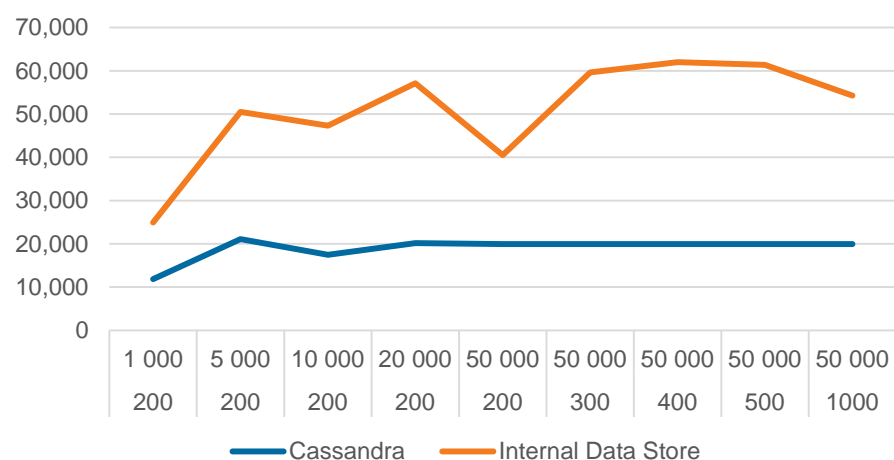


Data Store Performance

Creates/sec




Retrieves/sec



Thank you

Performance Results: Impact of Clustering on CDS



Performance Results: Lessons Learned

- Small transactions are dominated by initialization cost, charged to producer
 - Affects a single thread's ability to saturate backend
- Batching effectiveness goes up with backend latency
 - Many listeners, complex transaction processing, messaging latency
- Listening across shards results in heavy backend contention
 - Increases latency in notification delivery
 - Results in more events being batched hence spikes are observable
- Dynamic sharding improves performance with multiple applications
 - Per-application shards result in better isolation and improved parallelism
 - Single-threaded applications are unable to fully saturate IMDS

Getting & Running the Performance Test Code

- Clone coretutorials:

```
git clone https://git.opendaylight.org/gerrit/coretutorials.git
```

- Build a distribution with sharding performance tests:

```
cd coretutorials/clustering/shardingsimple  
mvn clean install
```

- Run the distribution:

```
./karaf/target/assembly/bin/karaf -Xmx4096m
```

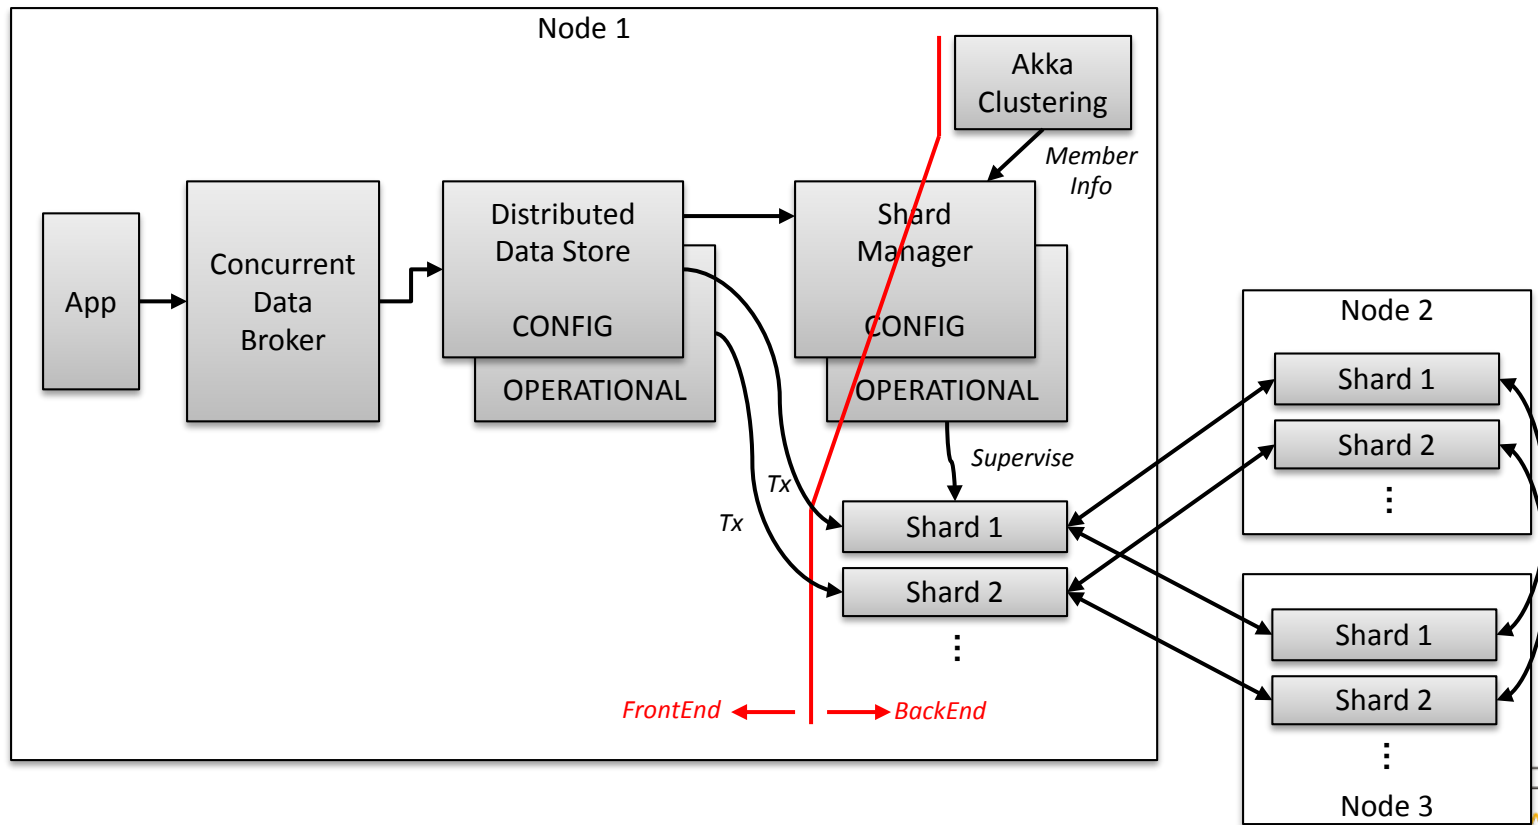
- Run the test script:

```
cd coretutorials/clustering/scripts/sharding  
./shardingbenchmark.py --help (will print all the parameters in the script)  
More info in coretutorials/clustering/scripts/sharding/site/asciidoc/scripts-user-manual.adoc
```

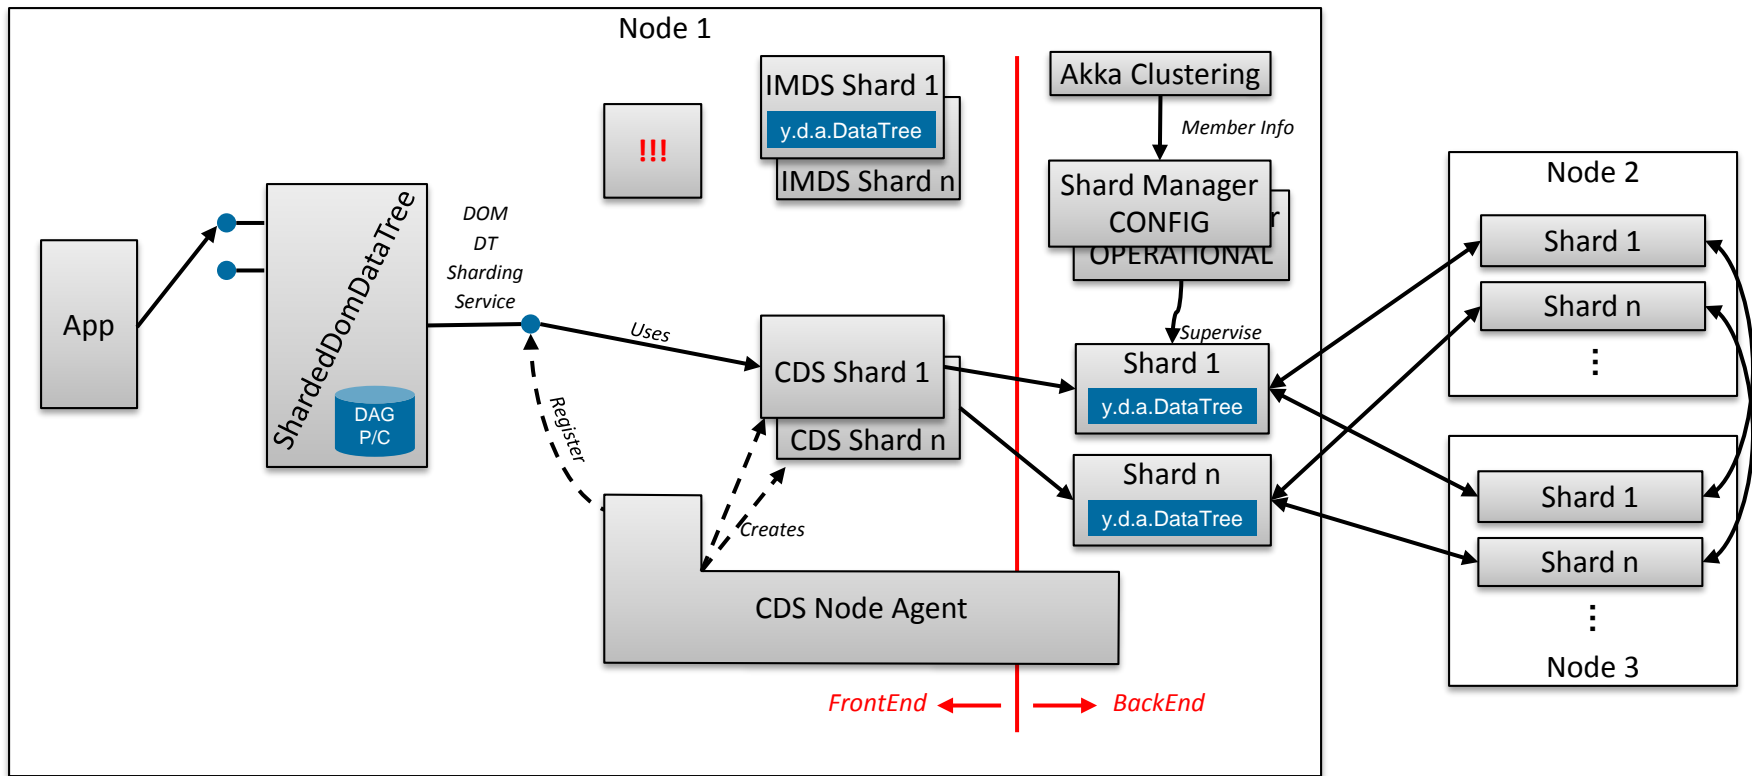
ODL Yang Resources

- YangTools main page:
 - https://wiki.opendaylight.org/view/YANG_Tools:Main
- Code Generation demo
 - https://wiki.opendaylight.org/view/Yang_Tools:Code_Generation_Demo
- Java “Binding Specification”:
 - https://wiki.opendaylight.org/view/YANG_Tools:YANG_to_Java_Mapping
- DLUX
 - Main page: https://wiki.opendaylight.org/view/OpenDaylight_dlux:Main
 - YangUI: https://wiki.opendaylight.org/view/OpenDaylight_dlux:yangUI-user
- Controller:
 - Swagger UI Explorer:
 - <http://localhost:8181/apidoc/explorer/index.html>
 - DLUX (Yangman):
 - <http://localhost:8181/dlux/index.html>

Existing Clustered Data Store - Details



Data Store Evolution in Boron and Beyond



y.d.a. = Yang Data API

DAB/PC – Directed Acyclic Graph Producer/Consumer